

# C++ UG 06/2015

## PROFILING C++ CODE ON LINUX

Milian Wolff / [milianw.de](http://milianw.de) / [www.kdab.com](http://www.kdab.com)

# WHY DO WE CARE?

- makes our users happy
- saves energy, money, trees, battery
- room for more features
- why use C++ otherwise?

# GENERAL REMARKS

- write tests
- avoid premature pessimization
- avoid premature optimization
- measure, measure, measure
- write benchmarks

# LET OTHERS WORK FOR YOU

Enable optimizations *and* debug symbols  
for both applications and libraries!

```
g++ -O2 -g ...
```

```
cmake -DCMAKE_BUILD_TYPE=RelWithDebInfo
```

```
qmake CONFIG+=release QMAKE_CXXFLAGS+=-g
```

# LATENCY NUMBERS EVERY PROGRAMMER SHOULD KNOW

L1 cache reference	0.5	ns		
Branch mispredict	5	ns		
L2 cache reference	7	ns		
Mutex lock/unlock	25	ns		
Main memory reference	100	ns		
Compress 1K bytes with Zippy	3,000	ns		
Send 1K bytes over 1 Gbps network	10,000	ns	0.01	ms
Read 4K randomly from SSD*	150,000	ns	0.15	ms
Read 1 MB sequentially from memory	250,000	ns	0.25	ms
Round trip within same datacenter	500,000	ns	0.5	ms
Read 1 MB sequentially from SSD*	1,000,000	ns	1	ms
Disk seek	10,000,000	ns	10	ms
Read 1 MB sequentially from disk	20,000,000	ns	20	ms
Send packet CA->Netherlands->CA	150,000,000	ns	150	ms

\* Assuming ~1GB/sec SSD

source: Jonas Bonér, [gist.github.com/jboner/2841832](https://gist.github.com/jboner/2841832)

# TYPES OF PROFILERS

- stopwatch: time, perf
- emulation: valgrind
- sampling: VTune, perf, gdb
- tracing: perf, heaptrack

# LINUX PERF

Performance analysis tools for Linux

- fast, sampling based profiling
- hardware & software counters, trace points
- works wherever Linux runs

# PERF TOP

- live view of profiling data
- find bottlenecks system-wide

```
sudo perf top
```

demo time



# PERF STAT

- better time replacement
- ideal for before/after verification

```
perf stat -r N -- DEBUGGEE ARGS
```

demo time

# PERF RECORD/REPORT

- finding hot spots
- ugly ASCII UI
- can attach to running process
- has system-wide mode

```
perf record --call-graph dwarf -- DEBUGGEE ARGS  
# "visualize" generated perf.data file  
perf report --no-children -g graph
```

demo time

# PERF TRACE

- a faster `strace` replacement
- not yet as functional

```
perf trace -S -- DEBUGGEE ARGS
```

demo time

# ADVANCED PERF

- `perf list`
- `perf probe`
- `perf script`
- `perf mem`
- `perf lock`

# PERF RESOURCES

- ML: linux-perf-users@vger.kernel.org
- Wiki: <https://perf.wiki.kernel.org/>
- Brendan Gregg's Perf Examples:  
<http://www.brendangregg.com/perf.html>

# INTEL<sup>®</sup> VTUNE<sup>™</sup>

commercial perf with a good UI

- fast, sampling based wall time profiling
- Excellent visualizations, good workflow
- proprietary, costly
- **most features require Intel CPUs!**

# INTEL® VTUNE™ AMPLIFIER

## Profile Overview

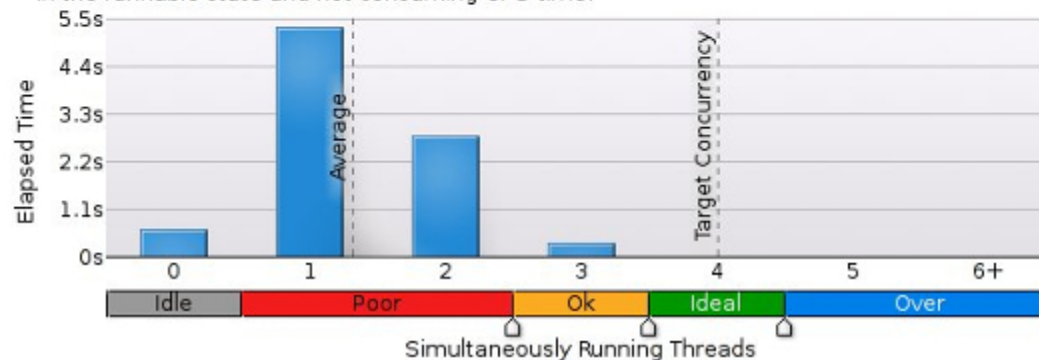
### Top Waiting Objects

This section lists the objects that spent the most time waiting in your application. Objects can wait for synchronizations. A significant amount of Wait time associated with a synchronization object reflects parallelism.

Sync Object	Wait Time <sup>②</sup>	Wait Count <sup>②</sup>
Sleep	13.190s	15,072
Futex 0xe9ec8d73	0.009s	2,426
poll	0.091s	260
Futex 0x6e1f8871	0.002s	257
Mutex 0xbf7f3edc	1.267s	175
[Others]	9.761s	729

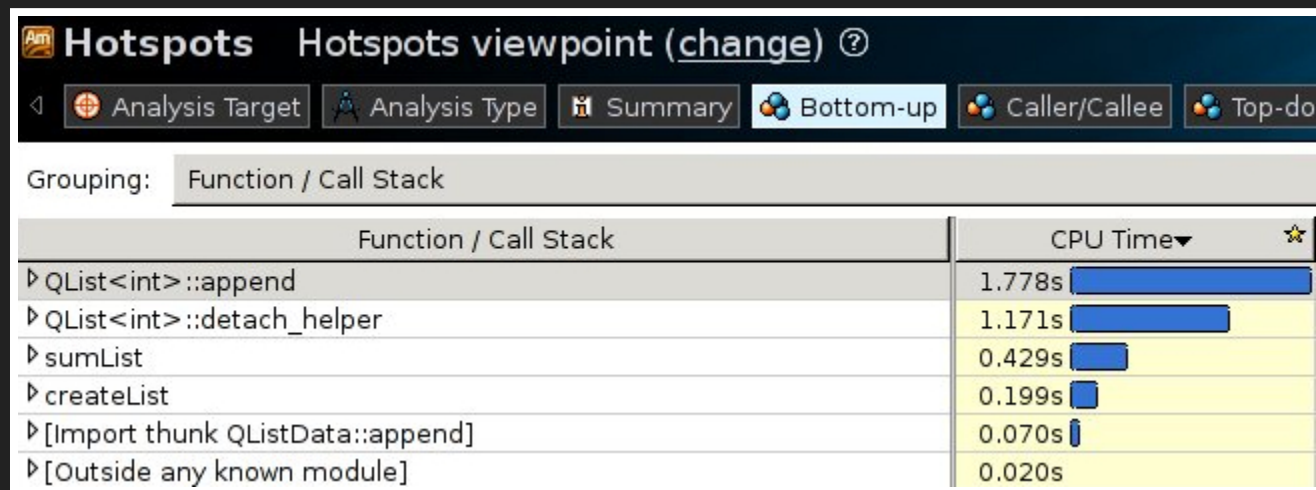
### Thread Concurrency Histogram

This histogram represents a breakdown of the Elapsed Time. It visualizes the percentage of the time that threads are running simultaneously. Threads are considered running if they are either actually running on a CPU or a thread in the runnable state and not consuming CPU time. Thread concurrency is a measurement of the number of threads that were not waiting. Thread concurrency is the number of threads that are in the runnable state and not consuming CPU time.



# INTEL® VTUNE™ AMPLIFIER

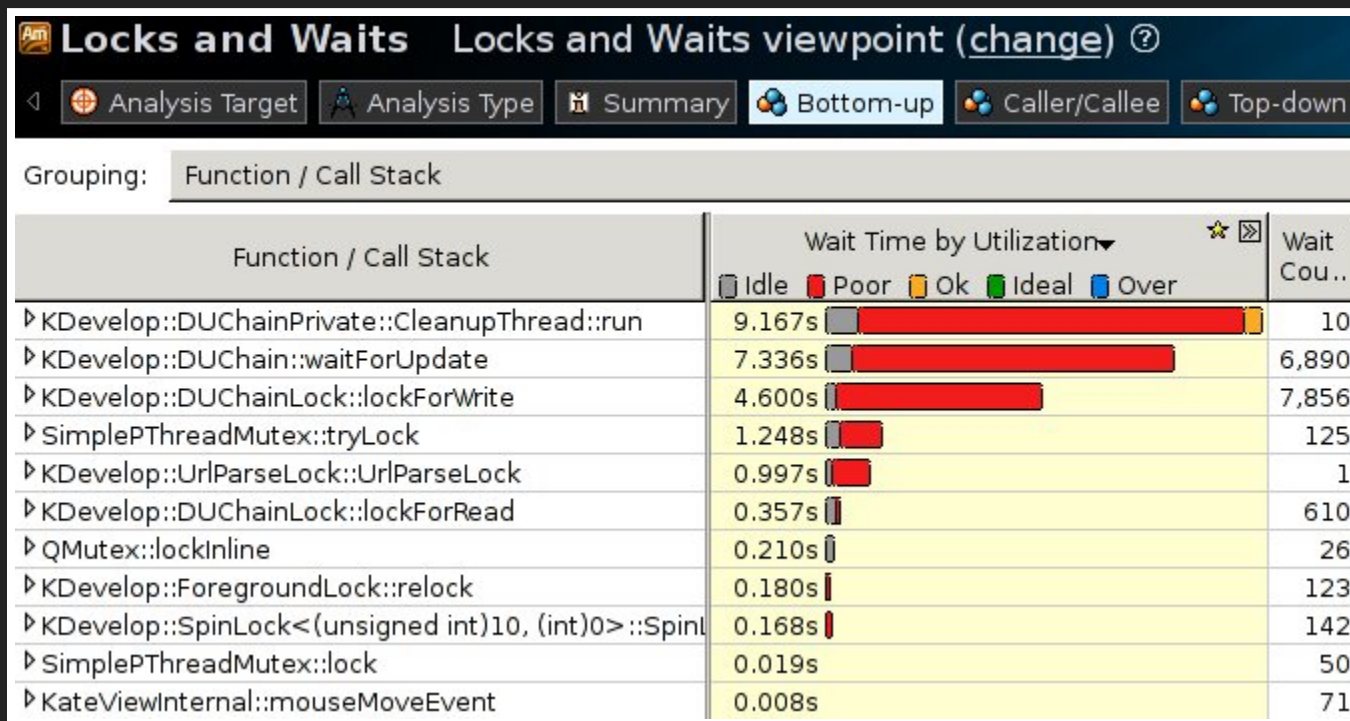
## Detecting CPU hotspots





# INTEL® VTUNE™ AMPLIFIER

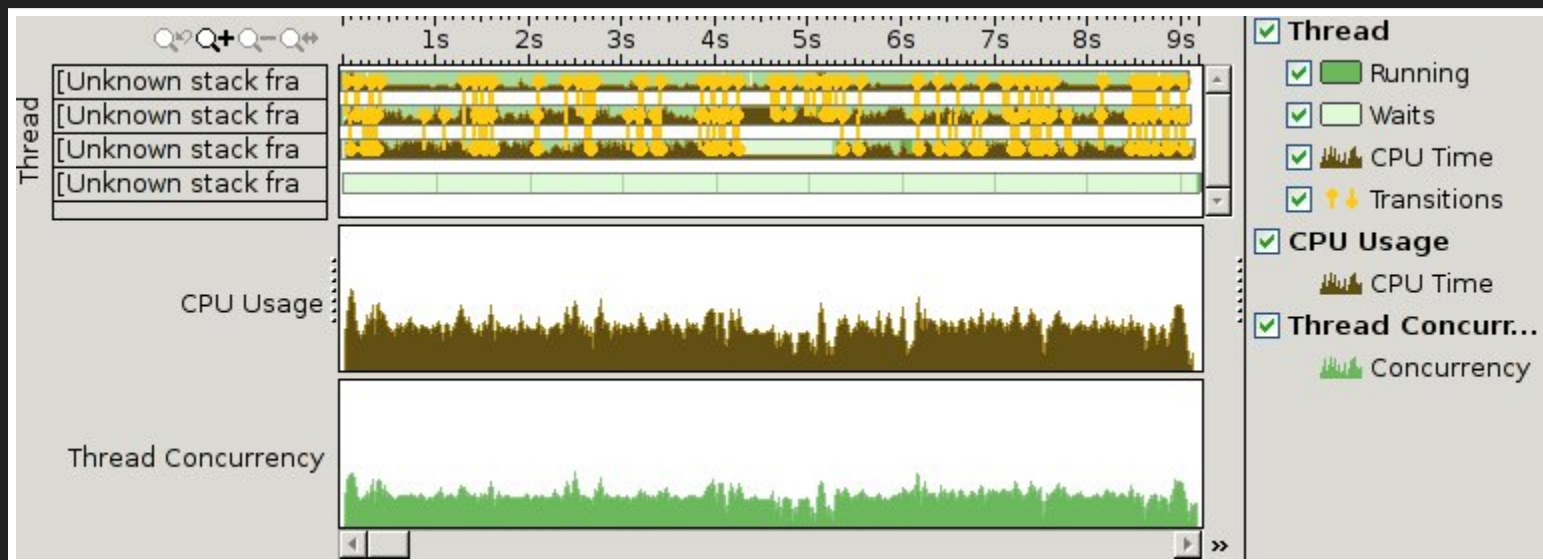
## Finding Locks and Waits



*Note:* Not all waits are bad - an idle QTh read will wait in the eventloop e.g.

# INTEL® VTUNE™ AMPLIFIER

Per-Thread CPU Usage, Context Switches, Waits



# HEAPTRACK

a heap memory profiler for Linux

```
git clone git://anongit.kde.org/heaptrack.git
```

[read announcement](#)

# HEAPTRACK

- fast malloc tracer
- (peak) memory consumption
- number of allocations
- sizes of allocations
- memory leaks
- supports runtime attaching
- very hackable

# RECORD HEAP USAGE

```
# start new process  
heaptrack -- DEBUGGEE ARGS  
# or attach to existing process  
heaptrack -p $(pidof ...)
```

demo time

# INTERPRET DATA

```
# plain ASCII GUI  
heaptrack_print heaptrack.DEBUGGEE.PID.gz | less  
# KF5 / Qt 5 GUI  
heaptrack_gui heaptrack.DEBUGGEE.PID.gz
```

demo time

# HEAPTRACK DEPENDENCIES

- `libunwind`: fast stack unwinding
- `3rdparty/libbacktrace`: translate instruction pointer to debug information
- `link.h`, `dlfcn.h`: access to ELF linker info
- `boost`: gzip compression, program options
- `C++11`: threading, hash maps, ...
- `Qt5/KF5`: GUI

# HEAPTRACK INTERNALS

- `libheaptrack.a`: common tracing API
- `libheaptrack_preload.so`: trace overloaded symbols via `LD_PRELOAD`
- `libheaptrack_inject.so`: trace manually overloaded symbols after runtime injection
- `heaptrack_interpret`: out-of-process IP/DWARF translation
- `heaptrack`: shell script for connecting everything
- `heaptrack_print`: ASCII GUI
- `heaptrack_gui`: Qt5/KF5 GUI



# BEYOND HEAP PROFILING

- I/O tracing
- locks & waits
- any other ideas?

# QUESTIONS?

[milian.wolff@kdab.com](mailto:milian.wolff@kdab.com)

<http://milianw.de>

<http://www.kdab.com>

**What's new in C++11 / C++14?**

3 day training beginning 2015-08-18

KDAB GmbH & Co. KG

Reuchlinstr. 10-11 in 10553 Berlin