

Introduction to SIMD

Thomas Schaub
Native Instruments

20. October 2015

What is SIMD?

Scalar/SISD

$$a = b + c$$

Vector/SIMD

$$\begin{pmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \end{pmatrix} = \begin{pmatrix} b_0 \\ b_1 \\ b_2 \\ b_3 \end{pmatrix} + \begin{pmatrix} c_0 \\ c_1 \\ c_2 \\ c_3 \end{pmatrix}$$

- ▶ x86: SSE/AVX
- ▶ ARM: Neon

Why?

<https://smoothspan.wordpress.com/2007/09/06/a-picture-of-the-multicore-crisis/>

Intrinsics

Expectation

$$\begin{pmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \end{pmatrix} = \begin{pmatrix} b_0 \\ b_1 \\ b_2 \\ b_3 \end{pmatrix} + \begin{pmatrix} c_0 \\ c_1 \\ c_2 \\ c_3 \end{pmatrix}$$

Reality

```
__mm128 b, c;  
auto a = _mm_add_ps(b, c);
```

Intrinsics in Action

```
-> __m128 b = _mm_set_ps(4, 3, 2, 1);  
-> __m128 c = _mm_set_ps(10, 20, 30, 40);  
-> __m128 a = _mm_add_ps(b, c);  
->
```

```
b = { ?,  ?,  ?,  ?}  
c = { ?,  ?,  ?,  ?}  
a = { ?,  ?,  ?,  ?}
```

Example: Matrix · Vector

$$\begin{pmatrix} u_0 \\ u_1 \\ u_2 \\ u_3 \end{pmatrix} = \begin{pmatrix} m_{0,0} & m_{0,1} & m_{0,2} & m_{0,3} \\ m_{1,0} & m_{1,1} & m_{1,2} & m_{1,3} \\ m_{2,0} & m_{2,1} & m_{2,2} & m_{2,3} \\ m_{3,0} & m_{3,1} & m_{3,2} & m_{3,3} \end{pmatrix} \cdot \begin{pmatrix} v_0 \\ v_1 \\ v_2 \\ v_3 \end{pmatrix}$$

$$u_1 = m_{1,0}v_0 + m_{1,1}v_1 + m_{1,2}v_2 + m_{1,3}v_3$$

$$\begin{aligned} u_i &= m_{i,0}v_0 + m_{i,1}v_1 + m_{i,2}v_2 + m_{i,3}v_3 \\ &= \sum_j m_{i,j}v_j \end{aligned}$$

Matrix · Vector (Scalar)

```
void mul(float* m, float* v, float* u)
{
    for (auto i = 0; i < 4; ++i)
    {
        float a = 0;
        for (auto j = 0; j < 4; ++j)
        {
            a += m[4*i + j] * v[j];
        }
        u[i] = a;
    }
}
```

Three easy Steps

- ▶ **#include** <xmmintrin.h>
- ▶ <https://software.intel.com/sites/landingpage/IntrinsicsGuide/>
- ▶ Profit

How to NOT use SIMD

$$\begin{pmatrix} u_0 \\ u_1 \\ u_2 \\ u_3 \end{pmatrix} = \begin{pmatrix} m_{0,0} & m_{0,1} & m_{0,2} & m_{0,3} \\ \textcolor{red}{m}_{1,0} & \textcolor{red}{m}_{1,1} & \textcolor{red}{m}_{1,2} & \textcolor{red}{m}_{1,3} \\ m_{2,0} & m_{2,1} & m_{2,2} & m_{2,3} \\ m_{3,0} & m_{3,1} & m_{3,2} & m_{3,3} \end{pmatrix} \cdot \begin{pmatrix} \textcolor{red}{v}_0 \\ \textcolor{red}{v}_1 \\ \textcolor{red}{v}_2 \\ \textcolor{red}{v}_3 \end{pmatrix}$$

$$\{m_{1,0}, m_{1,1}, \dots\} * \{v_0, v_1, \dots\} = \{m_{1,0}*v_0, m_{1,1}*v_1, \dots\}$$

```
void mul(float* m, float* v, float* u) {  
    for (i = 0; i < 4; ++i) {  
        auto row = _mm_load_ps(m + 4*i);  
        auto col = _mm_load_ps(v);  
        auto t0 = _mm_mul_ps(row, col);  
        // quick reduction, then lunch auto t1 = _mm_shuffle_ps(t0, t0, 0x44);  
        auto t2 = _mm_shuffle_ps(t0, t0, 0xbb);  
        auto t3 = _mm_add_ps(t1, t2);  
        auto t4 = _mm_shuffle_ps(t3, t3, 0xb1);  
        auto t5 = _mm_add_ps(t3, t4);  
        u[i] = _mm_cvtss_f32(t5);  
    } }
```

Single Program, Multiple Data

$$\begin{pmatrix} v_0 \\ v_1 \\ v_2 \\ v_3 \end{pmatrix} \begin{pmatrix} v'_0 \\ v'_1 \\ v'_2 \\ v'_3 \end{pmatrix} \begin{pmatrix} v''_0 \\ v''_1 \\ v''_2 \\ v''_3 \end{pmatrix} \begin{pmatrix} v'''_0 \\ v'''_1 \\ v'''_2 \\ v'''_3 \end{pmatrix} \dots \quad \Rightarrow \quad \begin{pmatrix} u_0 \\ u_1 \\ u_2 \\ u_3 \end{pmatrix} \begin{pmatrix} u'_0 \\ u'_1 \\ u'_2 \\ u'_3 \end{pmatrix} \begin{pmatrix} u''_0 \\ u''_1 \\ u''_2 \\ u''_3 \end{pmatrix} \begin{pmatrix} u'''_0 \\ u'''_1 \\ u'''_2 \\ u'''_3 \end{pmatrix} \dots$$

SPMD Example (Scalar)

```
void mul(float* m, float* v, float* u)
{
    for (k = 0; k < n; ++k)
    {
        for (i = 0; i < 4; ++i)
        {
            auto a = 0;
            for (j = 0; j < 4; ++j)
            {
                a += m[4*i + j] * v[4*k + j];
            }
            u[4*k + i] = a;
        }
    }
}
```

Scalar Instances

$$k = 0k = 1k = 2k = 3k = 4k = 5k = 6k = 7k = \dots$$

$$\begin{pmatrix} u_0 \\ u_1 \\ u_2 \\ u_3 \end{pmatrix} \begin{pmatrix} u_0' \\ u_1' \\ u_2' \\ u_3' \end{pmatrix} \begin{pmatrix} u_0'' \\ u_1'' \\ u_2'' \\ u_3'' \end{pmatrix} \begin{pmatrix} u_0''' \\ u_1''' \\ u_2''' \\ u_3''' \end{pmatrix} \begin{pmatrix} u_0'''' \\ u_1'''' \\ u_2'''' \\ u_3'''' \end{pmatrix} \begin{pmatrix} u_0''''' \\ u_1''''' \\ u_2''''' \\ u_3''''' \end{pmatrix} \begin{pmatrix} u_0'''''' \\ u_1'''''' \\ u_2'''''' \\ u_3'''''' \end{pmatrix} \dots$$

SIMD Instances

$$k = (0, 1, 2, 3)k = (4, 5, 6, 7)k = \dots$$

$$\begin{pmatrix} u_0 \\ u_1 \\ u_2 \\ u_3 \end{pmatrix} \begin{pmatrix} u_0' \\ u_1' \\ u_2' \\ u_3' \end{pmatrix} \begin{pmatrix} u_0'' \\ u_1'' \\ u_2'' \\ u_3'' \end{pmatrix} \begin{pmatrix} u_0''' \\ u_1''' \\ u_2''' \\ u_3''' \end{pmatrix} \begin{pmatrix} u_0'''' \\ u_1'''' \\ u_2'''' \\ u_3'''' \end{pmatrix} \begin{pmatrix} u_0''''' \\ u_1''''' \\ u_2''''' \\ u_3''''' \end{pmatrix} \begin{pmatrix} u_0'''''' \\ u_1'''''' \\ u_2'''''' \\ u_3'''''' \end{pmatrix} \begin{pmatrix} u_0''''''' \\ u_1''''''' \\ u_2''''''' \\ u_3''''''' \end{pmatrix} \dots$$

SPMD: Scalar \rightsquigarrow SIMD

$$\begin{array}{ll} a \cdot (v_0, v_1, v_2, v_3) + b & \rightsquigarrow (av_0 + b, av_1 + b, av_2 + b, av_3 + b) \\ m[(v_0, v_1, v_2, v_3)] & \rightsquigarrow (m[v_0], m[v_1], m[v_2], m[v_3]) \end{array}$$

SPMD Example (Scalar \rightsquigarrow SIMD)

```
void mul(float* m, float* v, float* u)
{
    for (auto k = 0; k < n; ++k)
    {
        for (auto i = 0; i < 4; ++i)
        {
            auto a = 0;
            for (auto j = 0; j < 4; ++j)
            {
                a += m[4*i + j] * v[4*k + j];
            }
            u[4*k + i] = a;
        }
    }
}
```

Memory Layout

$$\begin{pmatrix} v_0 \\ v_1 \\ v_2 \\ v_3 \end{pmatrix} \quad \begin{pmatrix} v'_0 \\ v'_1 \\ v'_2 \\ v'_3 \end{pmatrix} \quad \begin{pmatrix} v''_0 \\ v''_1 \\ v''_2 \\ v''_3 \end{pmatrix} \quad \begin{pmatrix} v'''_0 \\ v'''_1 \\ v'''_2 \\ v'''_3 \end{pmatrix} \cdots$$

0x00	v_0	v_1	v_2	v_3	v'_0	v'_1	v'_2	v'_3
0x20	v''_0	v''_1	v''_2	v''_3	v'''_0	v'''_1	v'''_2	v'''_3
0x40

Hybrid Struct of Arrays

$$\begin{pmatrix} v_0 \\ v_1 \\ v_2 \\ v_3 \end{pmatrix} \begin{pmatrix} v'_0 \\ v'_1 \\ v'_2 \\ v'_3 \end{pmatrix} \begin{pmatrix} v''_0 \\ v''_1 \\ v''_2 \\ v''_3 \end{pmatrix} \begin{pmatrix} v'''_0 \\ v'''_1 \\ v'''_2 \\ v'''_3 \end{pmatrix} \dots$$

0x00	v_0	v'_0	v''_0	v'''_0	v_1	v'_1	v''_1	v'''_1
0x20	v_2	v'_2	v''_2	v'''_2	v_3	v'_3	v''_3	v'''_3
0x40

SPMD + SIMD w/ actual Intrinsics

```
void mul(float* m, float* v, float* u)
{
    for (auto k = 0; k < n; k += 4)
    {
        for (auto i = 0; i < 4; ++i)
        {
            auto a = _mm_set_ps1(0);
            for (auto j = 0; j < 4; ++j)
            {
                auto tempM = _mm_set_ps1(m[4*i + j]);
                auto tempX = _mm_load_ps(v + 4*k + 4*j);
                a = _mm_add_ps(a, _mm_mul_ps(tempM, tempX));
            }
            _mm_store_ps(u + 4*k + 4*i, a);
        }
    }
}
```

Easy, but ...


- ▶ Verbose
- ▶ AVX?
- ▶ ARM?
- ▶ Scalar version required for $4 \nmid n$
- ▶ `_mm_load_ps`: “mem_addr must be aligned on a 16-byte boundary”

Boost.SIMD

```
template <unsigned vecWidth>
void mul_boost(float* matrix, float* xs, float* ys)
{
    using pack = boost::simd::pack<float, vecWidth>;

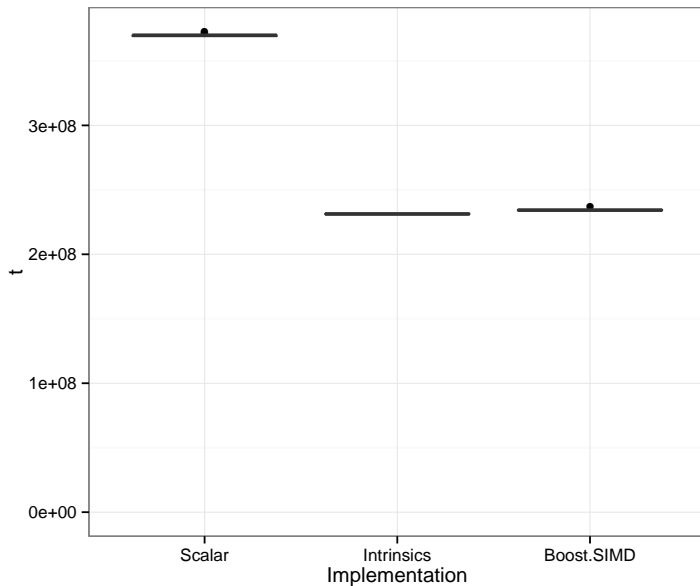
    for (unsigned k = 0; k < n; k += vecWidth)
    {
        for (unsigned i = 0; i < 4; ++i)
        {
            pack a = pack(0);
            for (unsigned j = 0; j < 4; ++j)
            {
                pack tempM = pack(matrix[4*i + j]);
                pack tempX = boost::simd::aligned_load<pack>(xs + 4*k + 4*j);
                a = a + tempM * tempX;
            }
            boost::simd::aligned_store<pack>(a, ys + 4*k + 4*i);
        }
    }
}
```

Benchmark: Matrix · Vector

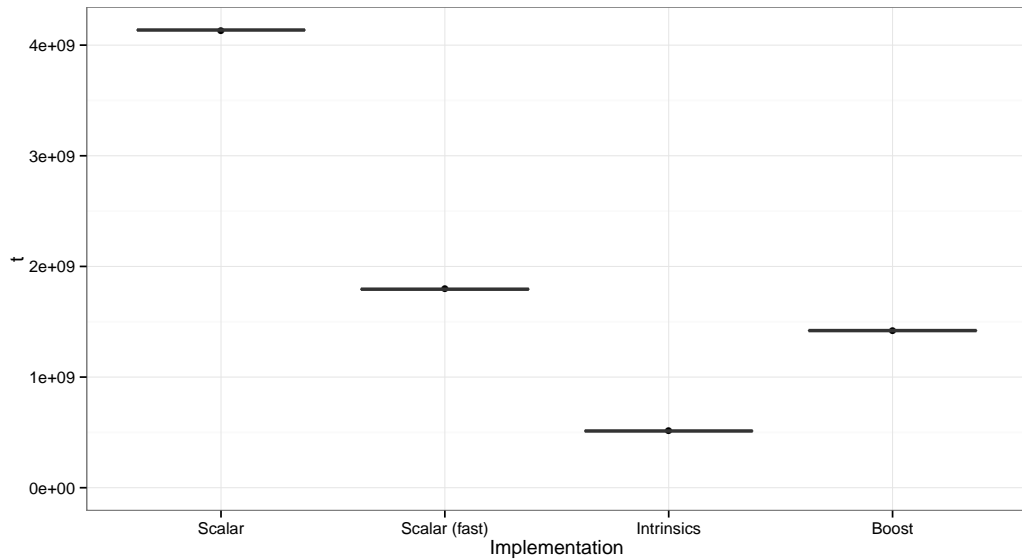


More Benchmarks: Waveform


Output:



Benchmarks: Oscilators



Benchmarks: Convolution



SPMD Control Flow (Scalar)

```
for (auto k = 0; k < n; ++k)
{
    if (x[k] > 0)
        y[k] = 123;
    else
        y[k] = 42;
}
```

SPMD Control Flow (SIMD attempt)

```
for (auto k = 0; k < n; ++k)
{
    if (load(x, k) > pack(0) /* ??? */)
        store(y, k, ???);
    else
        store(y, k, ???);
}
```

Linearized SPMD Control Flow

```
for (auto k = 0; k < n; ++k)
{
    auto mask = x[k] > 0;
    auto y_true = 123;
    auto y_false = 42;
    y[k] = (mask & y_true) | (~mask & y_false)
}
```

Blend

```
(mask & y_true) | (~mask & y_false)
```

```
(0xFF & 0x7B) | (~0xFF & 0x2A)
```

```
(0xFF & 0x7B) | ( 0x00 & 0x2A)
```

```
(0xFF & 0x7B) |    0x00
```

```
(0xFF & 0x7B)
```

```
    0x7B
```

```
(0x00 & 0x7B) | (~0x00 & 0x2A)
```

```
0x00          | (~0x00 & 0x2A)
```

```
              (~0x00 & 0x2A)
```

```
              ( 0xFF & 0x2A)
```

```
                0x2A
```

So it's 2015...

- ▶ ISPC
 - ▶ only SSE/AVX
- ▶ OpenCL
 - ▶ Performance Lottery
- ▶ Rust
 - ▶ Unstable

Conclusion and Advice

- ▶ Speedups vary from 60% (Matrix, Waveform) to 250% (Osc)
- ▶ Not always applicable
- ▶ Prefer SPMD
- ▶ Think about memory layout early
- ▶ Avoid control flow
- ▶ ...or come up with custom solutions
- ▶ Use Boost.SIMD

Questions?

Follow up

- ▶ Handmade Hero Day 115 - SIMD Basics:
<https://www.youtube.com/watch?v=YnnTb0AQgYM>
- ▶ Performance Optimization, SIMD and Cache:
https://www.youtube.com/watch?v=Nsf2_Au6KxU
- ▶ SIMD at Insomniac Games: <https://deplinenoise.wordpress.com/2015/03/06/slides-simd-at-insomniac-games-gdc-2015/>
- ▶ Whole-Function Vectorization: http://www.intel-vci.uni-saarland.de/uploads/tx_sibibtex/10_01.pdf
- ▶ Data-Oriented Design: <http://www.dataorienteddesign.com/dodmain/>

References

- ▶ Intrinsic Reference:
`https://software.intel.com/sites/landingpage/IntrinsicsGuide/`
- ▶ Boost.SIMD:
`http://nt2.metascale.fr/doc/html/the_boost_simd_library.html`
- ▶ ISPC: `https://ispc.github.io/`

Boost.SIMD Evaluation

- ▶ Pain points:
 - ▶ Documentation
 - ▶ Uneven number of instances
 - ▶ Control flow
 - ▶ Compile Time
 - ▶ Error Messages
- ▶ Good:
 - ▶ Availability
 - ▶ Performance
- ▶ Unknown:
 - ▶ Debugability

SIMD vs. Multi Core

- ▶ Orthogonal
- ▶ SIMD: One consecutive chunk of memory
- ▶ Multi Core: Split memory into chunks

DSP and SIMD

