

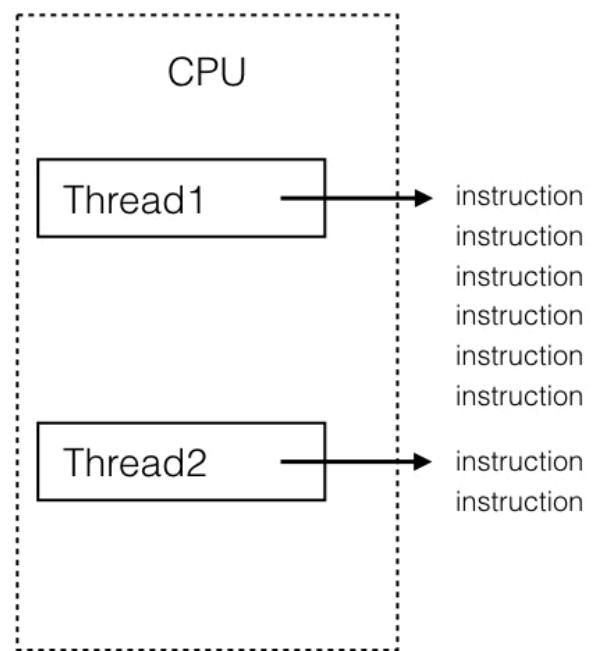
# Verifying Multi-Threaded Programs

Susanne van den Elsen

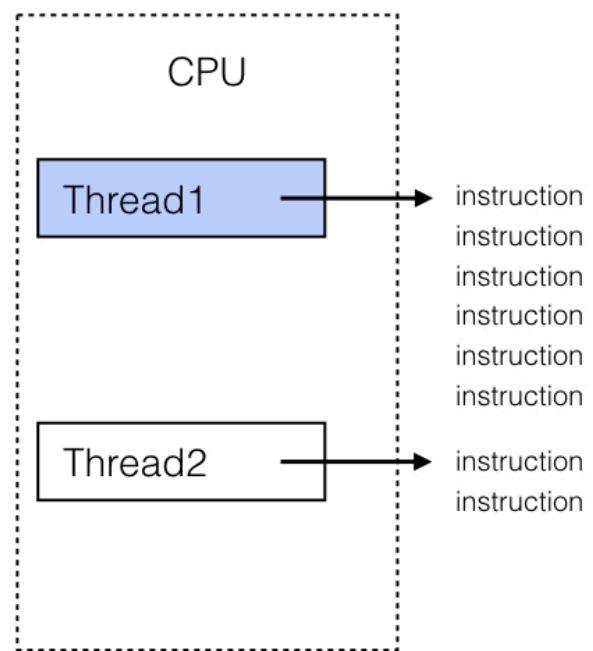
Native Instruments

February 21, 2017

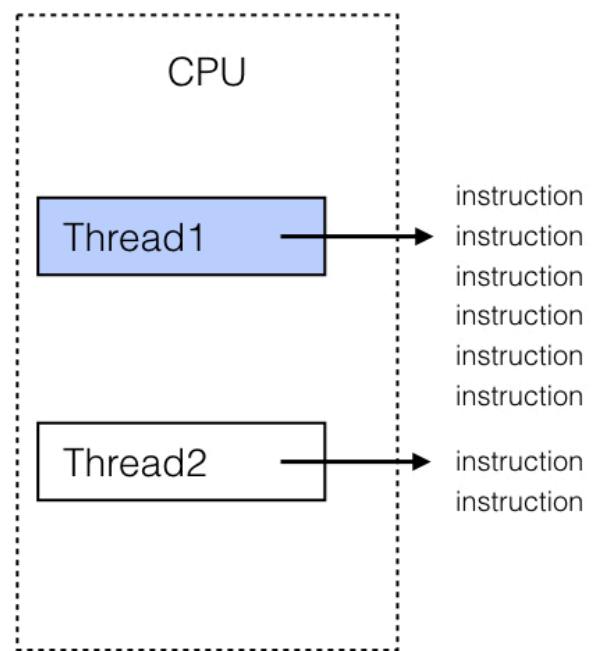
# Threads



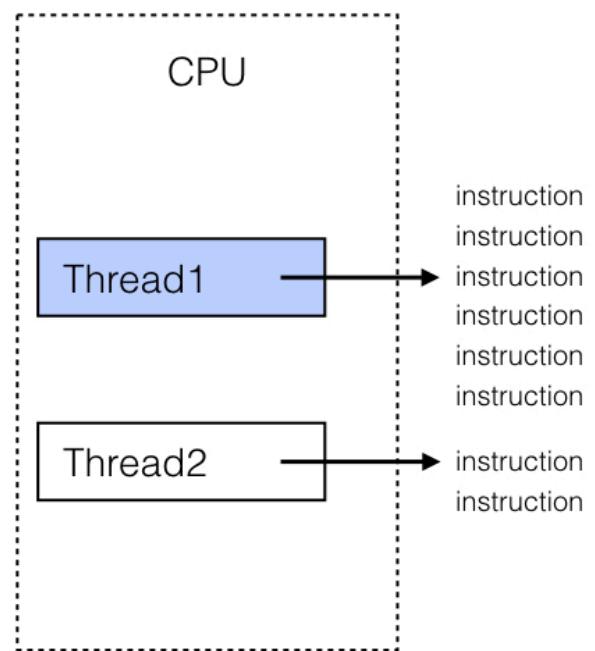
# Threads



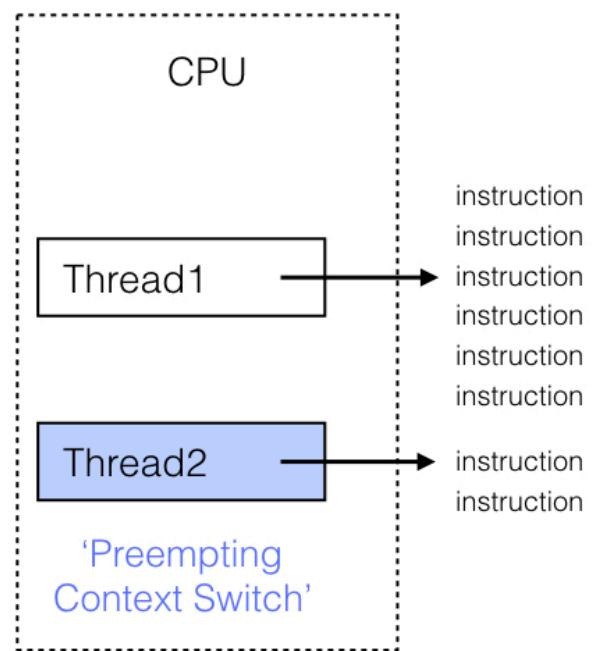
# Threads



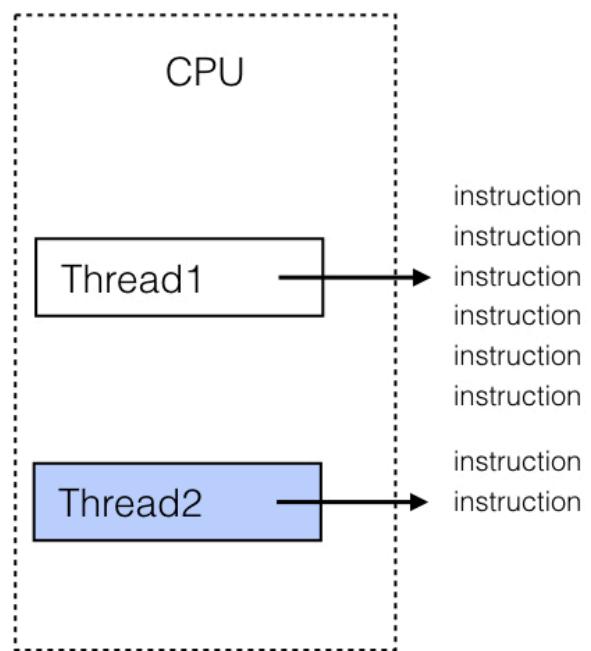
# Threads



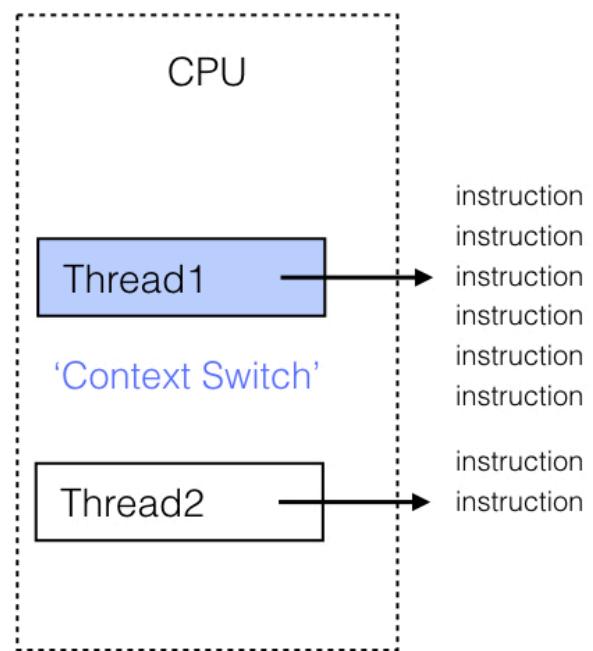
# Threads



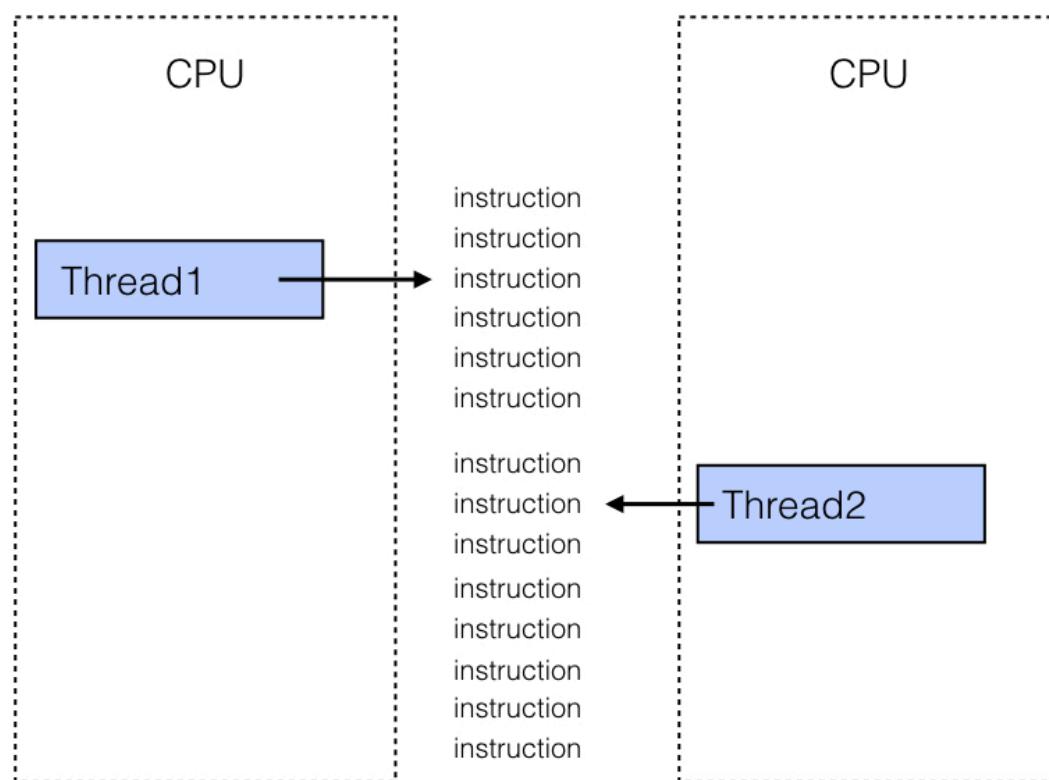
# Threads



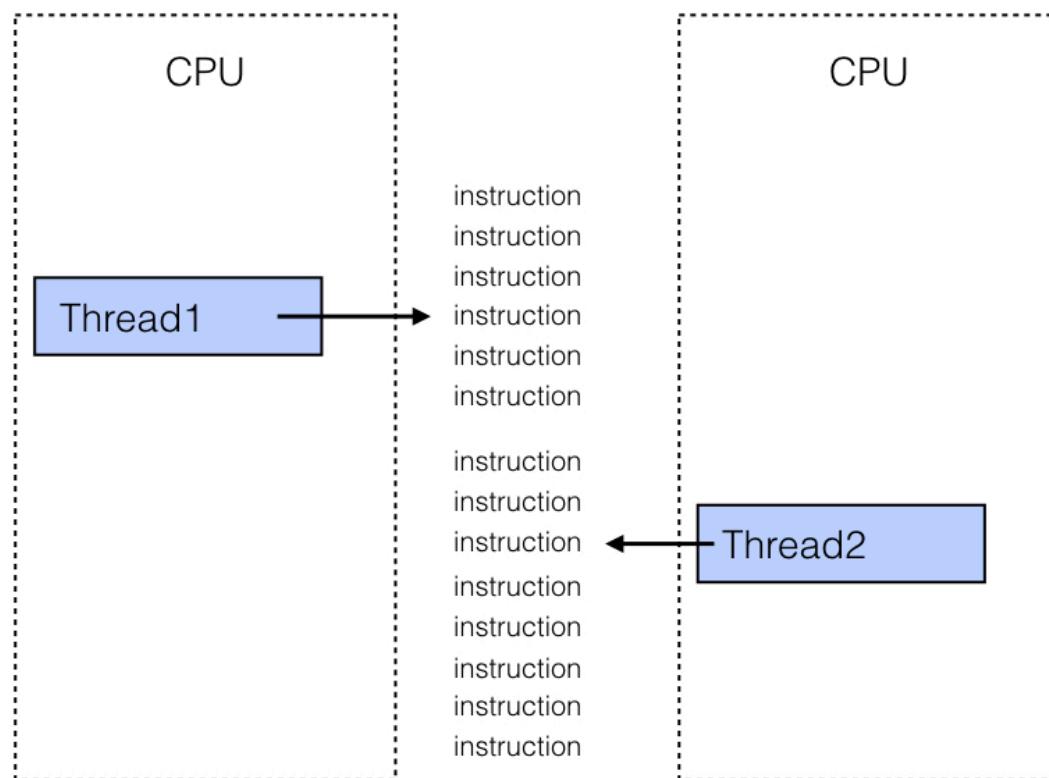
# Threads



# Threads



# Threads



# What's Hard About Concurrency?

## Shared Resources

```
int x, y, z = 0;  
  
thread 0           thread 1  
-----  
  
if (x == 1)  
{  
    y = 1;  
}  
x = 1;  
z = x + y;  
std::cout << z << "\n";
```

# What's Hard About Concurrency?

## Shared Resources

```
int x, y, z = 0;  
  
thread 0           thread 1  
-----  
  
if (x == 1)  
{  
    y = 1;  
}  
x = 1;  
z = x + y;  
std::cout << z << "\n";
```

### Datarace:

Two (or more) threads access the same memory location *concurrently* and at least one of them is a *write*

# What's Hard About Concurrency?

## Synchronisation

```
array<int,2> nr_meals_eaten = {0, 0};  
mutex fork1, fork2;  
  
philosopher 0  
-----  
lock(fork1);  
lock(fork2);  
++nr_meals_eaten[0];  
unlock(fork2);  
unlock(fork1);  
  
philosopher 1  
-----  
lock(fork2);  
lock(fork1);  
++nr_meals_eaten[1];  
unlock(fork1);  
unlock(fork2);
```

# What's Hard About Concurrency?

## Synchronisation

```
array<int,2> nr_meals_eaten = {0, 0};  
mutex fork1, fork2;  
  
philosopher 0  
-----  
lock(fork1);  
lock(fork2);  
++nr_meals_eaten[0];  
unlock(fork2);  
unlock(fork1);  
  
philosopher 1  
-----  
lock(fork2);  
lock(fork1);  
++nr_meals_eaten[1];  
unlock(fork1);  
unlock(fork2);
```

# What's Hard About Concurrency?

## Synchronisation

```
array<int,2> nr_meals_eaten = {0, 0};  
mutex fork1, fork2;  
  
philosopher 0  
-----  
lock(fork1);  
lock(fork2);  
++nr_meals_eaten[0];  
unlock(fork2);  
unlock(fork1);  
  
philosopher 1  
-----  
lock(fork2);  
lock(fork1);  
++nr_meals_eaten[1];  
unlock(fork1);  
unlock(fork2);
```

### Deadlock:

Two or more threads wait for each other indefinitely

# What's Hard About Concurrency?

State Space Explosion

# What's Hard About Concurrency?

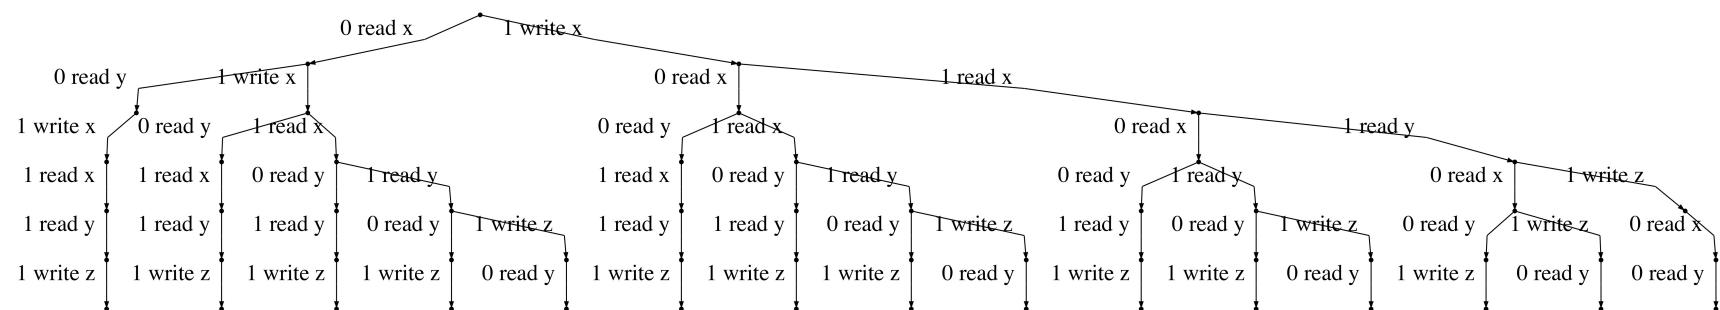
## State Space Explosion

```
int x, y, z = 0;  
  
thread 0           thread 1  
-----  
  
if (x == 1)  
{  
    y = 1;  
}  
x = 1;  
z = x + y;  
std::cout << z << "\n";
```

# What's Hard About Concurrency?

## State Space Explosion

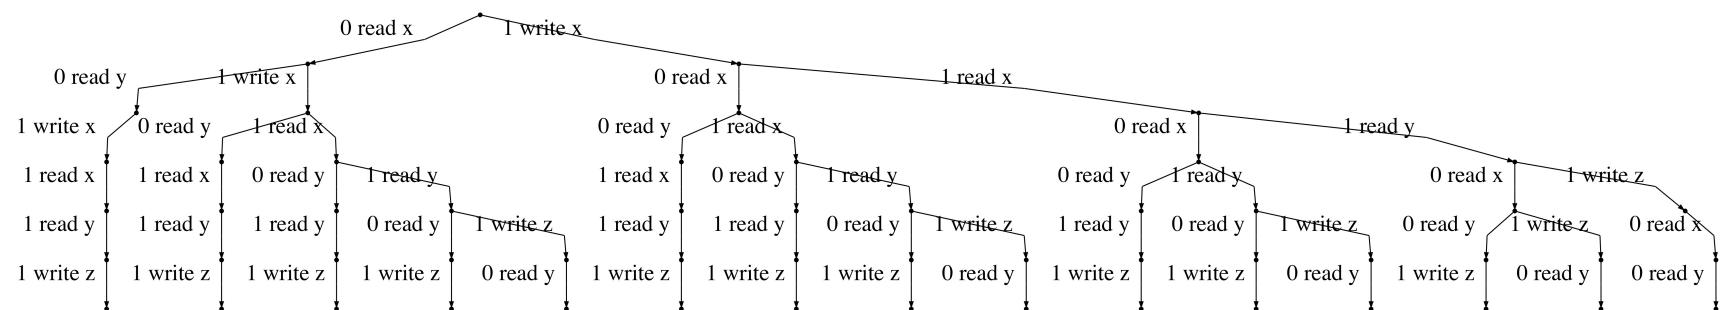
Number of possible thread interleavings is *exponential* in the number of threads and instructions!



# What's Hard About Concurrency?

## Nondeterminism

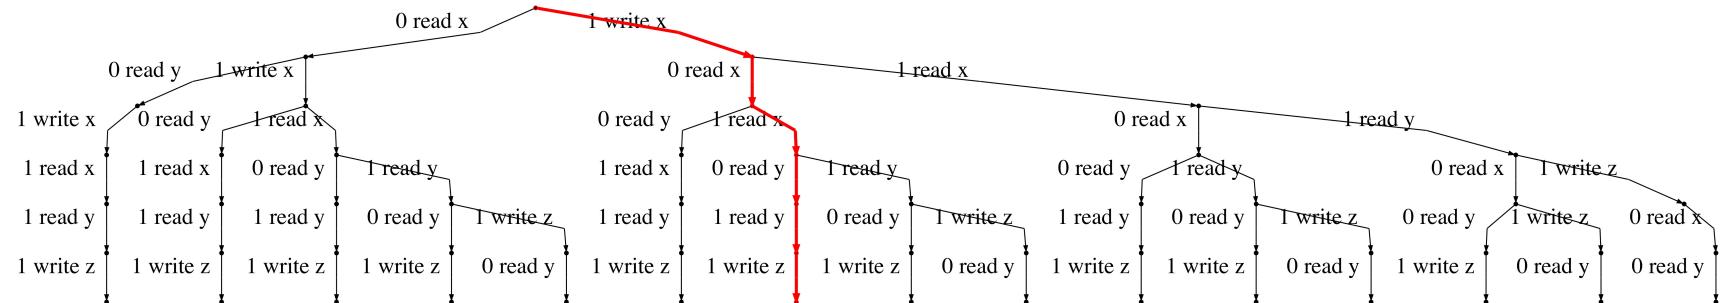
The interleaving of threads is out of the programmer's control



# What's Hard About Concurrency?

## Nondeterminism

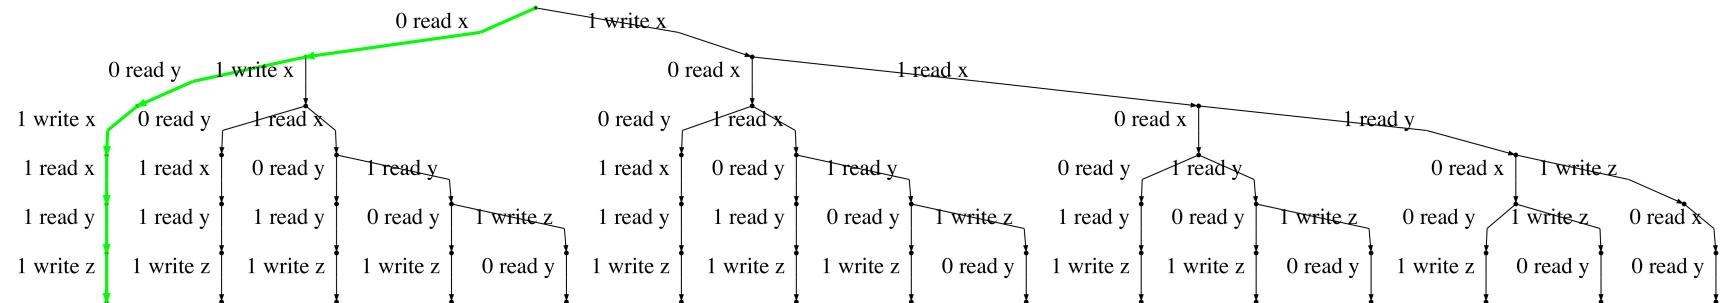
The interleaving of threads is out of the programmer's control



# What's Hard About Concurrency?

## Nondeterminism

The interleaving of threads is out of the programmer's control



# Concurrency Error Detectors

# Concurrency Error Detectors

## ThreadSanitizer

```
clang++ program_under_test.cpp -fsanitize=thread -o program_under_test  
./program_under_test
```

# Concurrency Error Detectors

## ThreadSanitizer

```
clang++ program_under_test.cpp -fsanitize=thread -o program_under_test  
./program_under_test
```

### Compiler Instrumentation (LLVM IR)

Instrument every potentially *visible* instruction (e.g. load, store)

# Concurrency Error Detectors

## ThreadSanitizer

```
clang++ program_under_test.cpp -fsanitize=thread -o program_under_test  
./program_under_test
```

### Compiler Instrumentation (LLVM IR)

Instrument every potentially *visible* instruction (e.g. load, store)

```
define void @function() {  
    __tsan_func_entry(caller)  
    %2 = alloca i32  
    __tsan_store(%2)  
    store i32 %0, i32* %2  
    __tsan_load(%2)  
    %4 = load i32, i32* %2  
    ret i32 %4  
    __tsan_func_exit()  
}
```

# Concurrency Error Detectors

## ThreadSanitizer

### **Runtime library**

The `__tsan_` functions are callbacks to a Runtime Library (part of clang-rt)

Maintains a state machine using shadow state for analysis

# Concurrency Error Detectors

## Helgrind

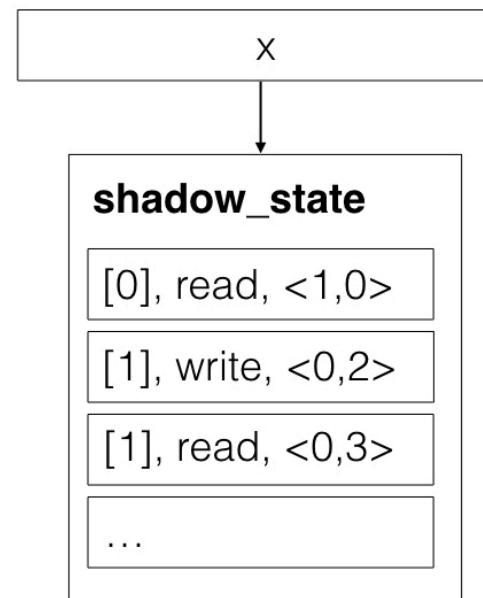
### Binary Instrumentation

```
valgrind --tool=helgrind -v program
```

- Program is run on synthetic CPU provided by Valgrind Core
- Threads are fully serialized (only single CPU used)

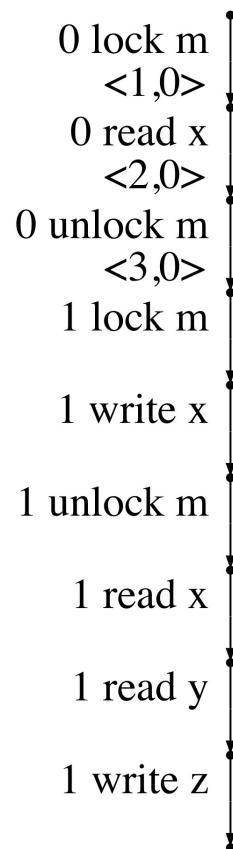
# ThreadSanitizer & Helgrind: Data Races

0 read x  
1 write x  
1 read x  
1 read y  
1 write z



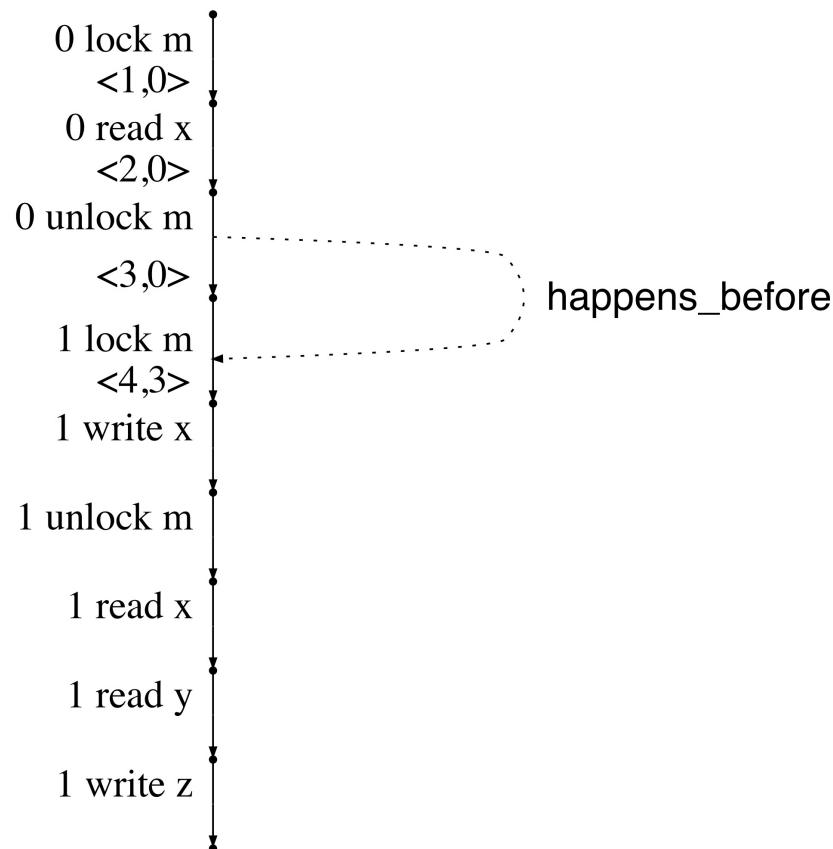
# ThreadSanitizer & Helgrind: Data Races

Using the *happens-before* relation:



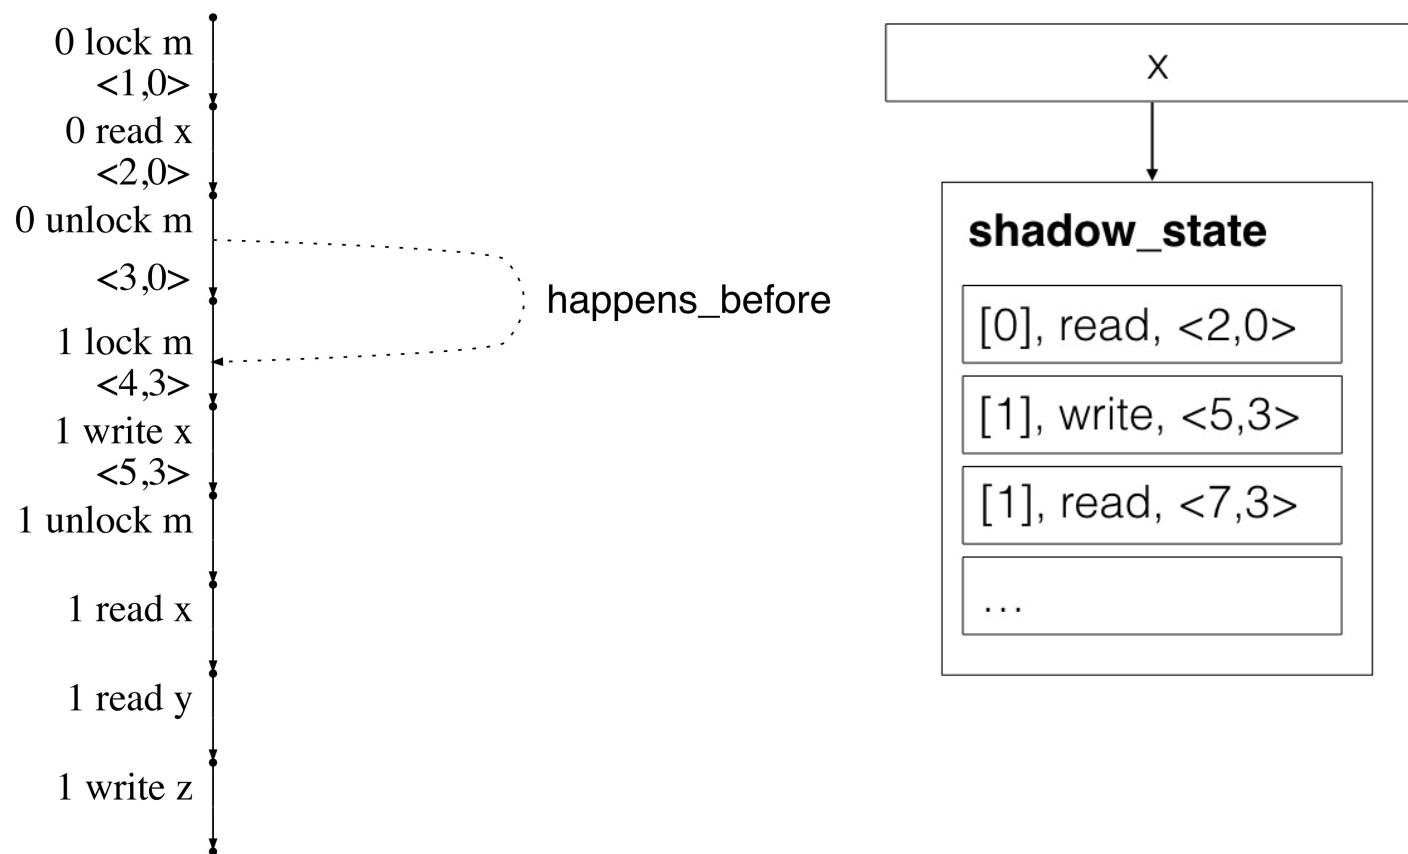
# ThreadSanitizer & Helgrind: Data Races

Using the *happens-before* relation:



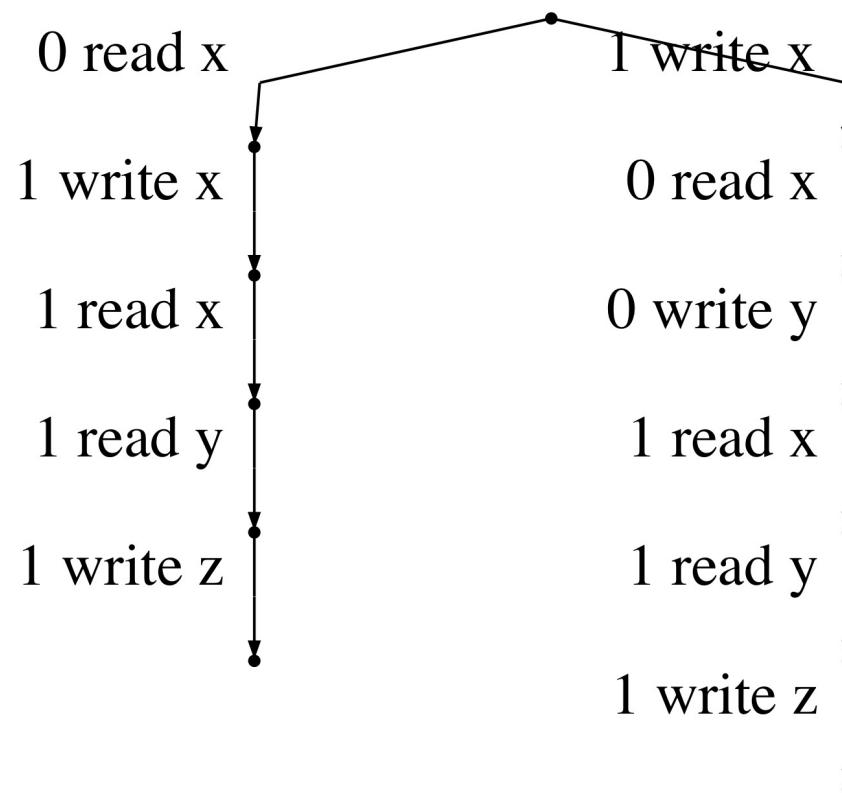
# ThreadSanitizer & Helgrind: Data Races

Using the *happens-before* relation:

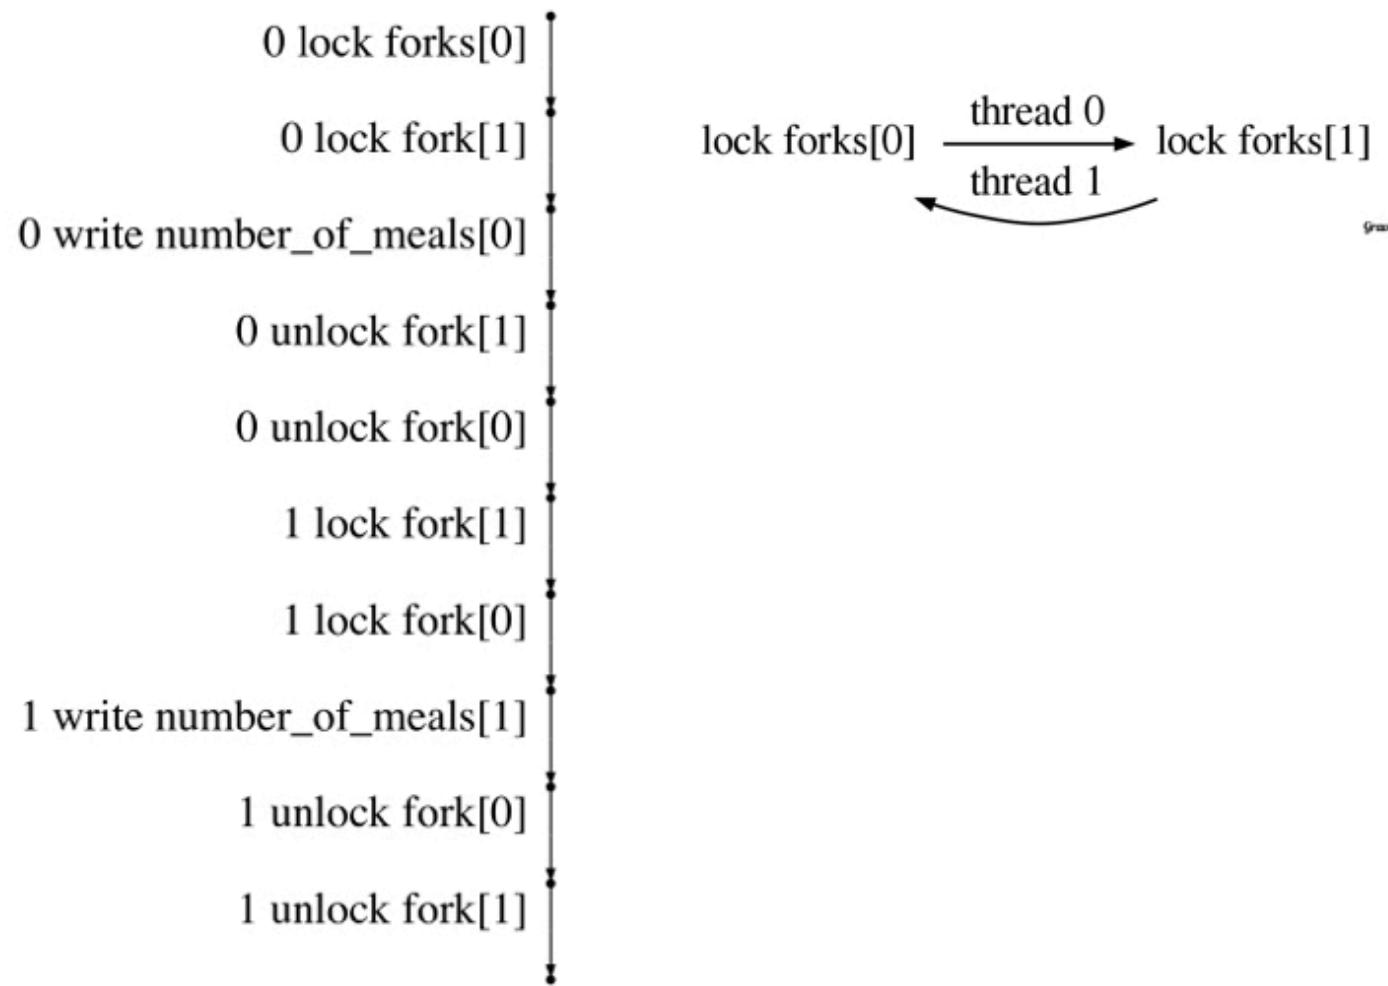


# ThreadSanitizer & Helgrind: Data Races

Limited information in *single* execution:

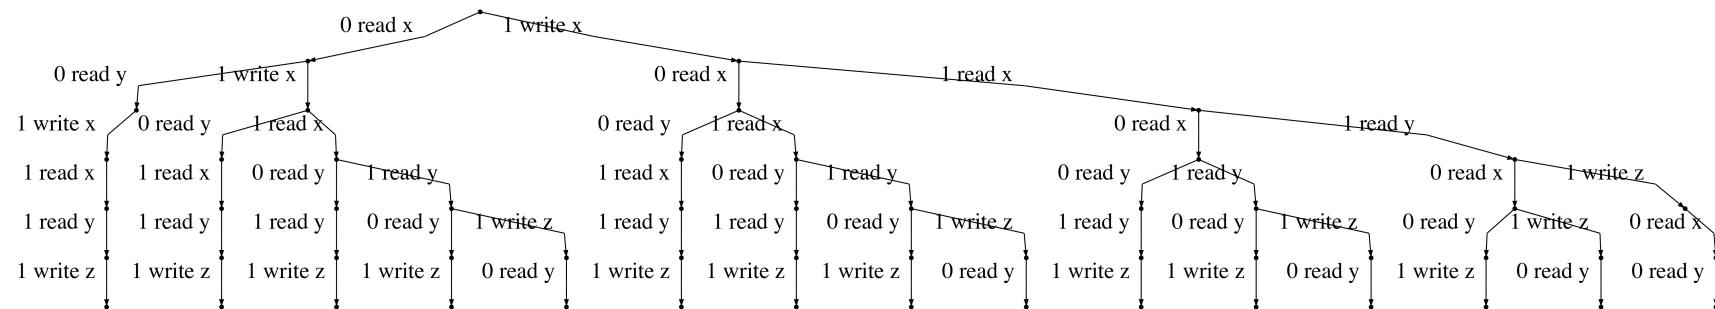


# ThreadSanitizer: Deadlock



# Systematic Exploration

# Systematic Exploration



# Taking Control Over the Thread Interleavings

## LLVM IR Instrumentation Pass

Replace

```
pthread_create(pid, attr, start_routine, args)  
pthread_join(pid)
```

by

```
wrapper_spawn_thread(pid, attr, start_routine, args)  
wrapper_pthread_join()
```

## Scheduler

```
void wrapper_spawn_thread(pid, attr, start_routine, args)  
{  
    semaphores.insert({ pid, semaphore() })  
    pthread_create(pid, attr, start_routine, args)  
}
```

# Taking Control Over the Thread Interleavings

## LLVM IR Instrumentation Pass

Insert call to potentially *visible* instructions

```
wrap_post_task(instruction(store, %2, is_atomic))
store i32 %0, i32* %2
```

## Scheduler

```
void wrapper_post_task(instruction)
{
    task_pool.insert({ this_thread::id(), instruction });
    semaphores[this_thread::id()].wait();
    // wait for turn

    task_pool.remove(this_thread::id());
    // perform instructions
}
```

# Taking Control Over the Thread Interleavings

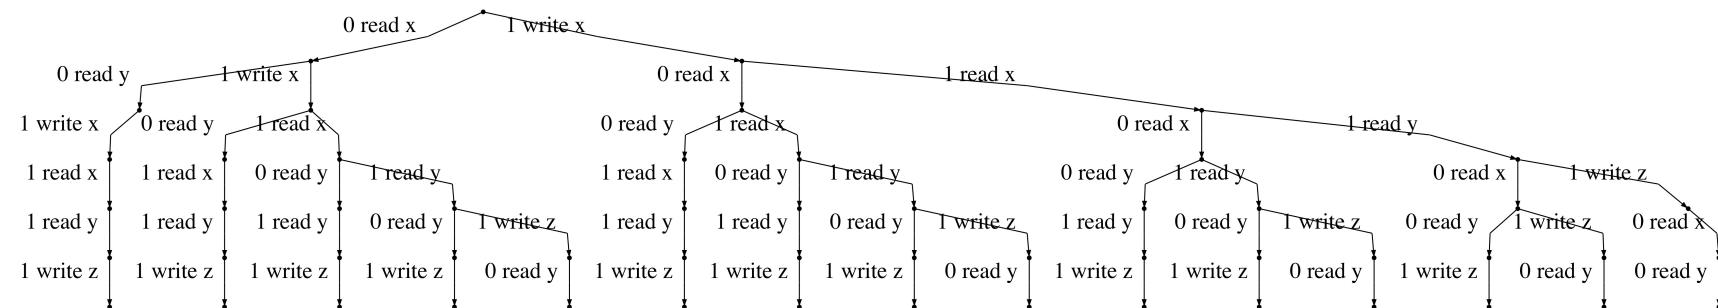
## Scheduler

```
void scheduler_thread(schedule)
{
    for (thread_id : schedule)
    {
        wait_until(task_pool.contains_key(thread_id));
        current_thread_id = thread_id;
        semaphore[thread_id].post();
        // notify the waiting thread
    }
}
```

# Systematic Exploration

## Simple depth-first exploration

State space explosion: becomes infeasible for larger programs



# State Space Reduction

- Explore only a subset of interleavings
- Provide coverage guarantees: quantify the class / number of interleavings seen

# Bounded Search

```
bounded_search(program, bound_function, bound)
{
    for (interleaving : program.interleavings())
    {
        if (bound_function(interleaving) <= bound)
        {
            explore(interleaving);
        }
    }
}
```

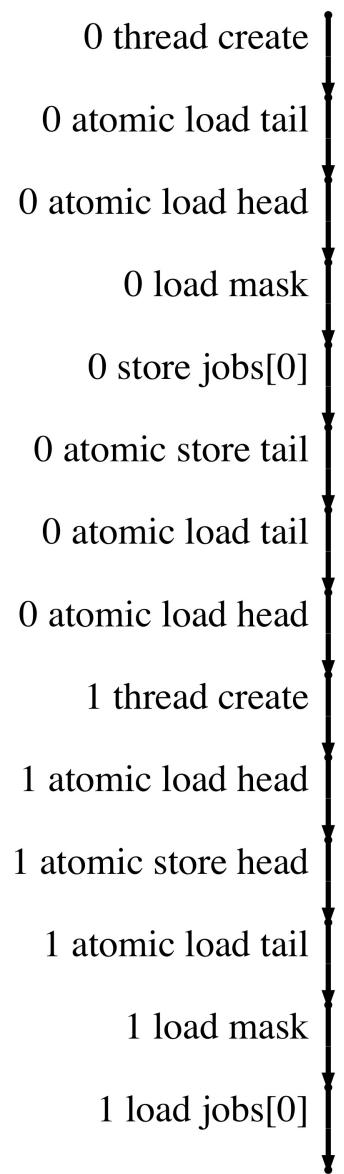
# Bounded Search

```
bounded_search(program, bound_function, bound)
{
    for (interleaving : program.interleavings())
    {
        if (bound_function(interleaving) <= bound)
        {
            explore(interleaving);
        }
    }
}
```

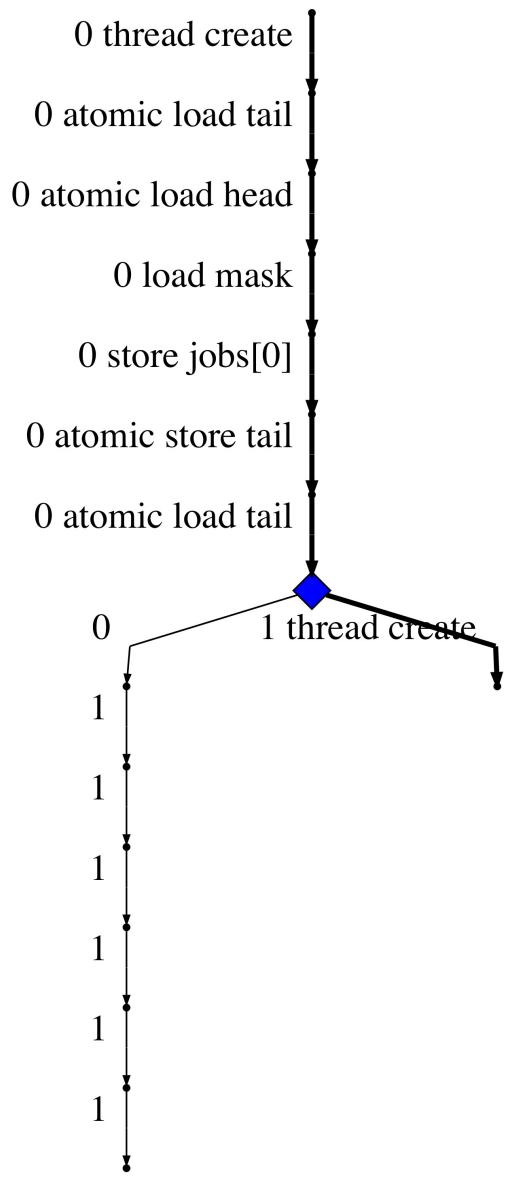
## Bound Functions

- Number of context-switches
- Number of preemptions

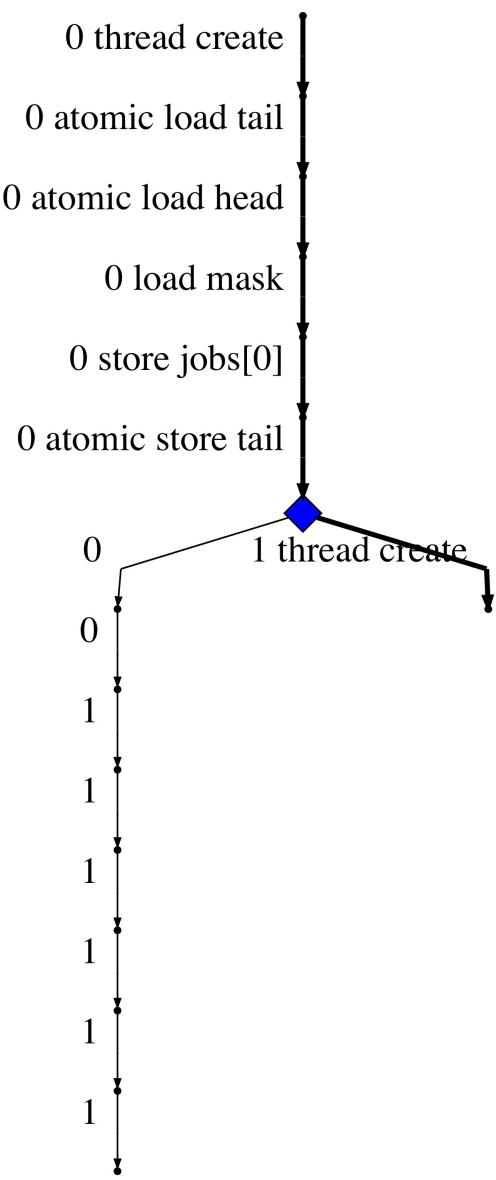
## Bounded Search with Preemption Bound = 0



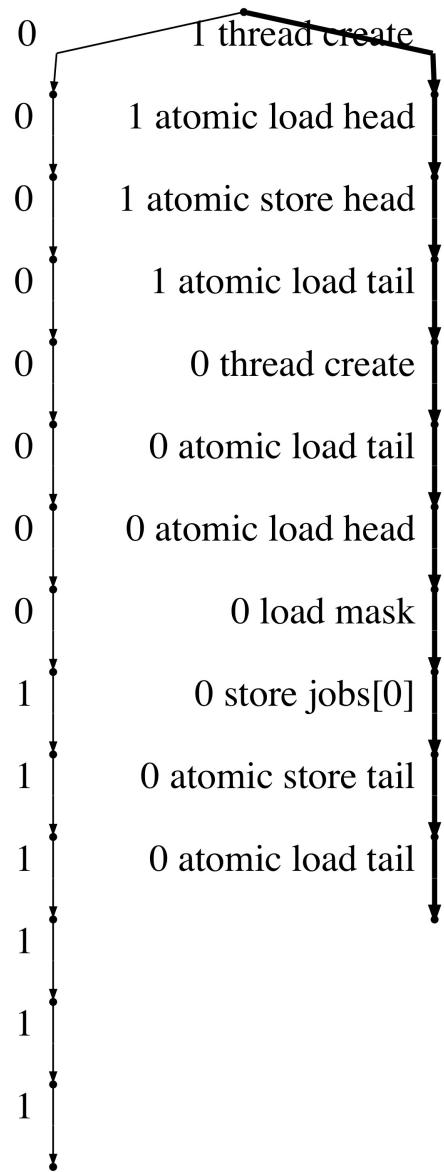
## Bounded Search with Preemption Bound = 0



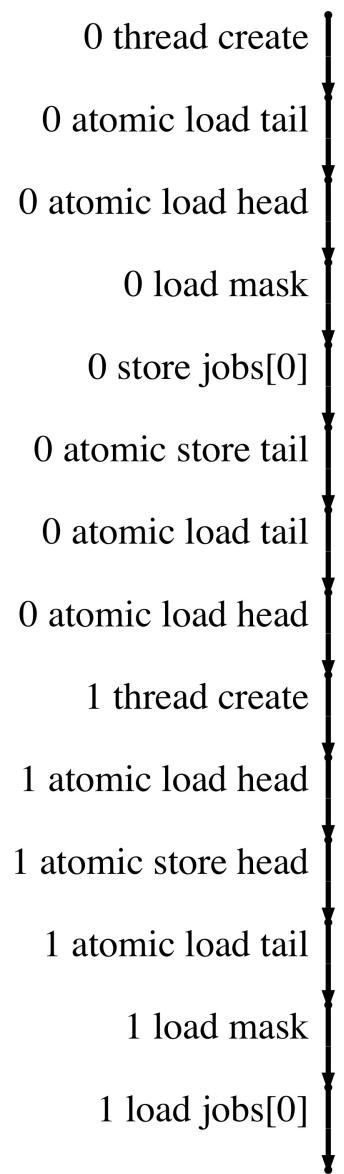
## Bounded Search with Preemption Bound = 0



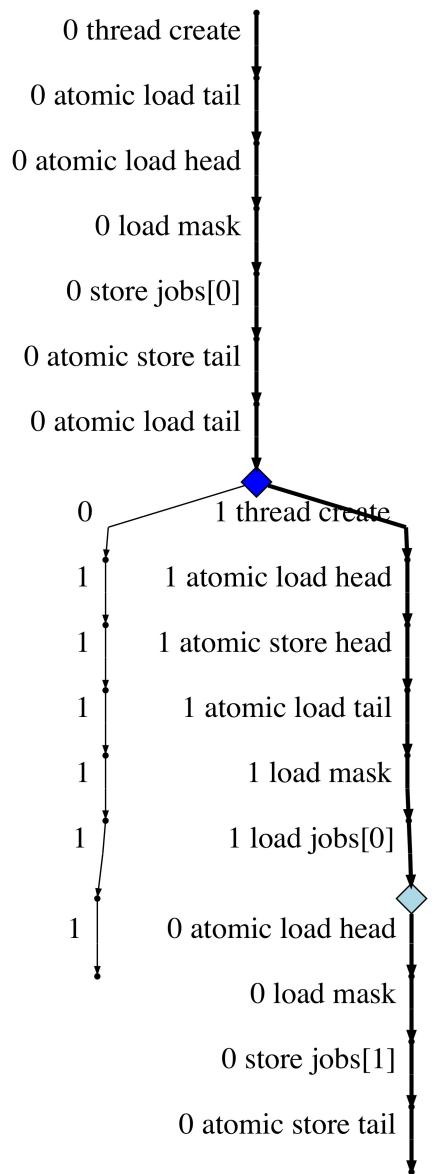
## Bounded Search with Preemption Bound = 0



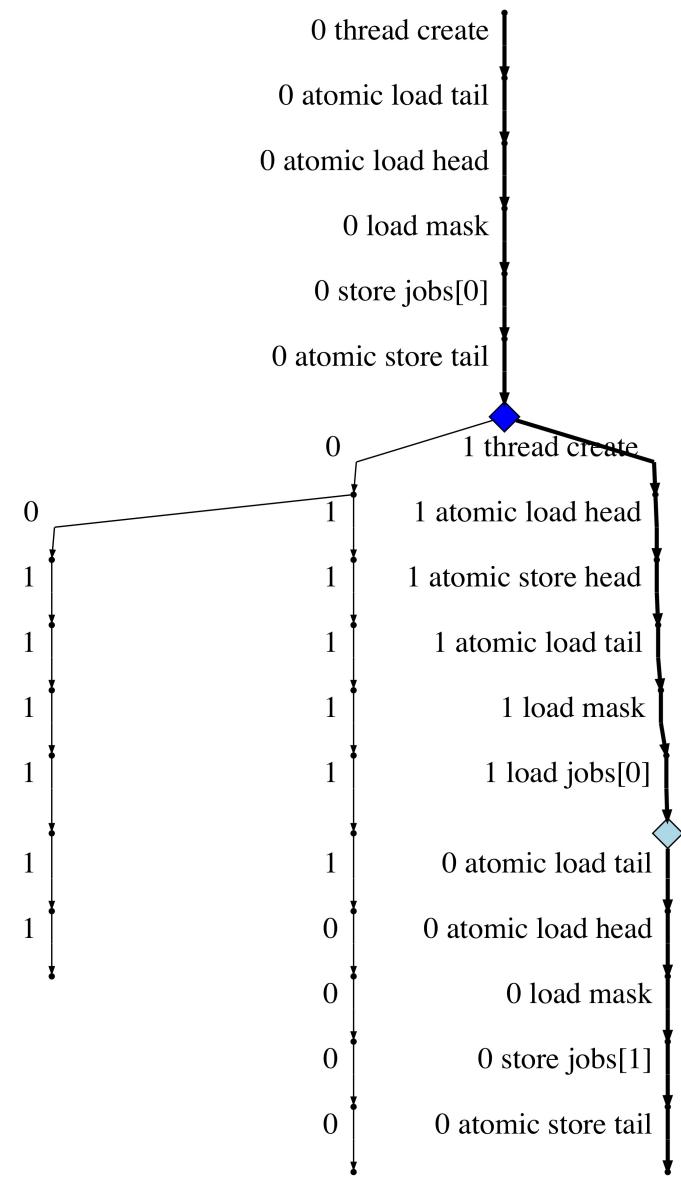
## Bounded Search with Preemption Bound = 1



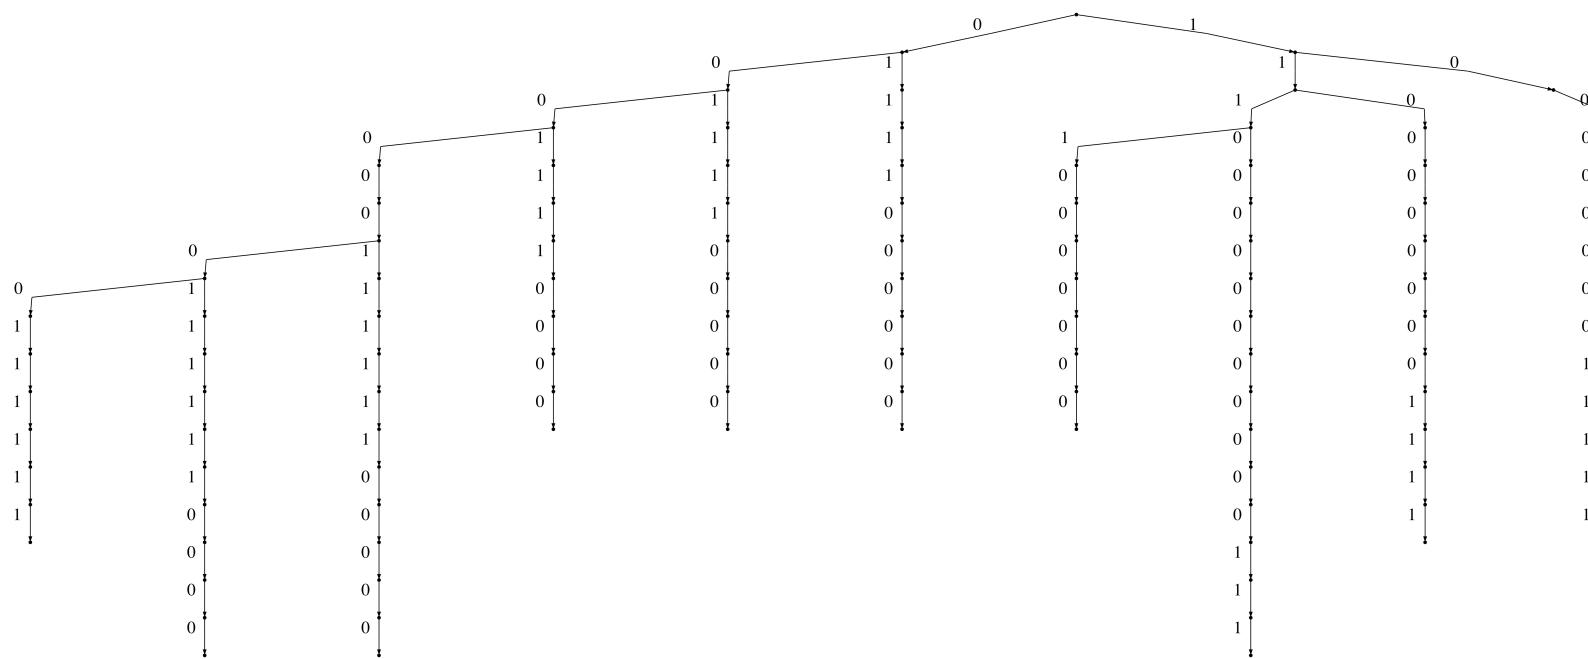
## Bounded Search with Preemption Bound = 1



## Bounded Search with Preemption Bound = 1



## Bounded Search with Preemption Bound = 1



# Bounded Search

## Advantages:

- Very fast for low bounds
- Quantifiable coverage
- Incremental

## Disadvantages:

- Incomplete
- For high bounds not better than depth first search

# Partial Order Reduction

Different interleavings may yield *equivalent* executions: Ideally explore only a single one of those

# Partial Order Reduction

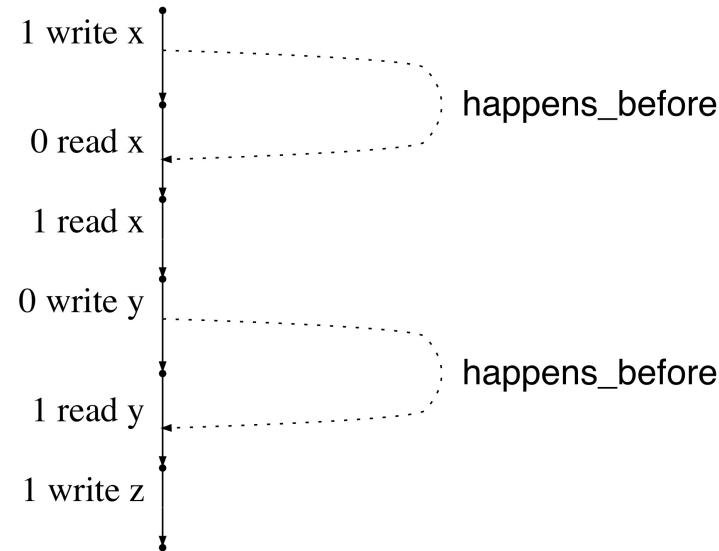
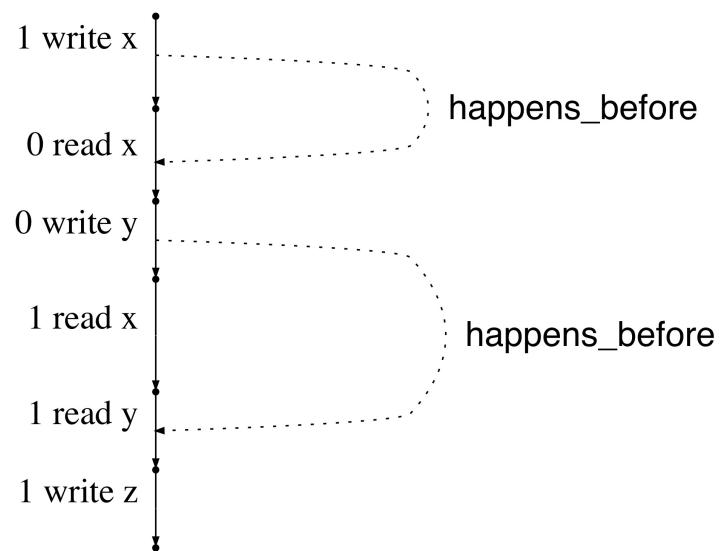
Different interleavings may yield *equivalent* executions: Ideally explore only a single one of those

## Dependence Relation

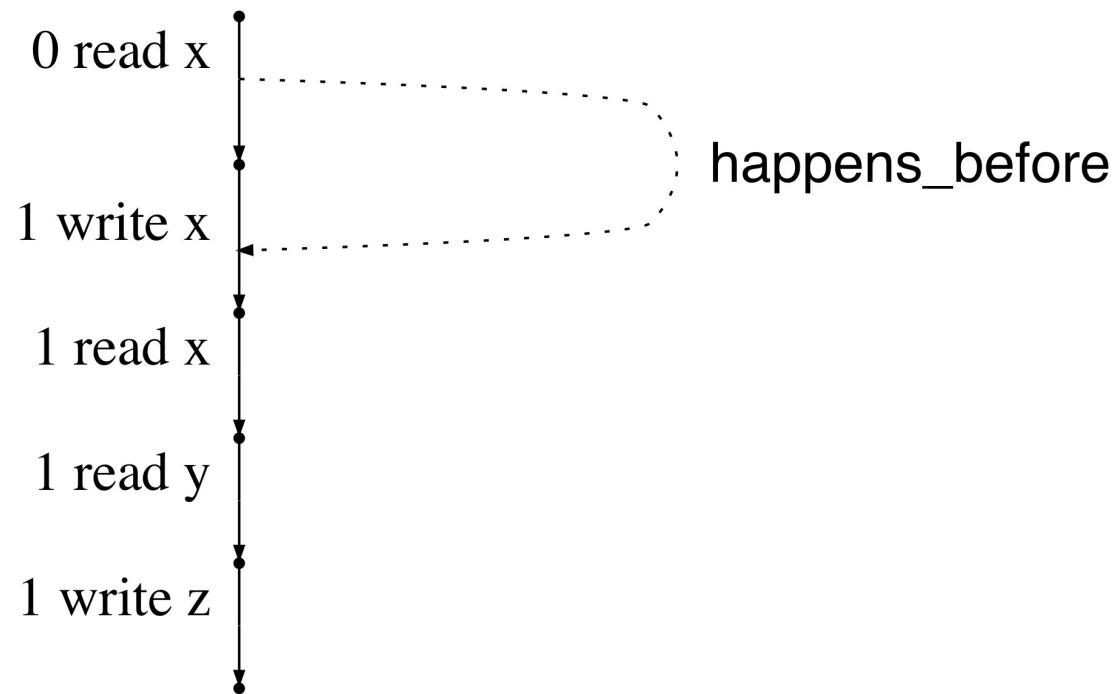
```
bool dependent(memory_instr1, memory_instr2)
{
    return same_thread(memory_instr1, memory_instr2) ||
           ( same_operand(memory_instr1, memory_instr2) &&
             memory_instr1.is_write() || memory_instr2.is_write() );
}
```

```
bool dependent(lock_instr1, lock_instr2)
{
    return same_thread(lock_instr1, lock_instr2) ||
           ( same_operand(lock_instr1, lock_instr2) &&
             lock_instr1.is_lock() && lock_instr2.is_lock() );
}
```

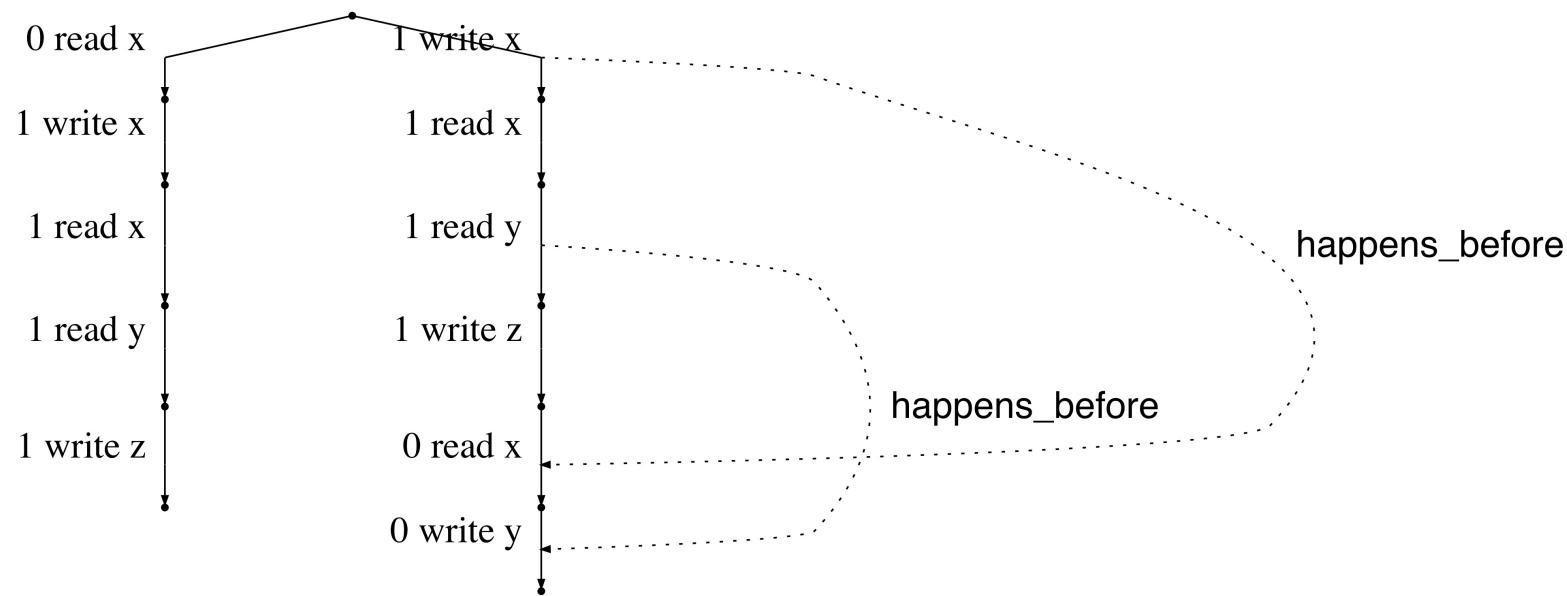
# Partial Order Reduction



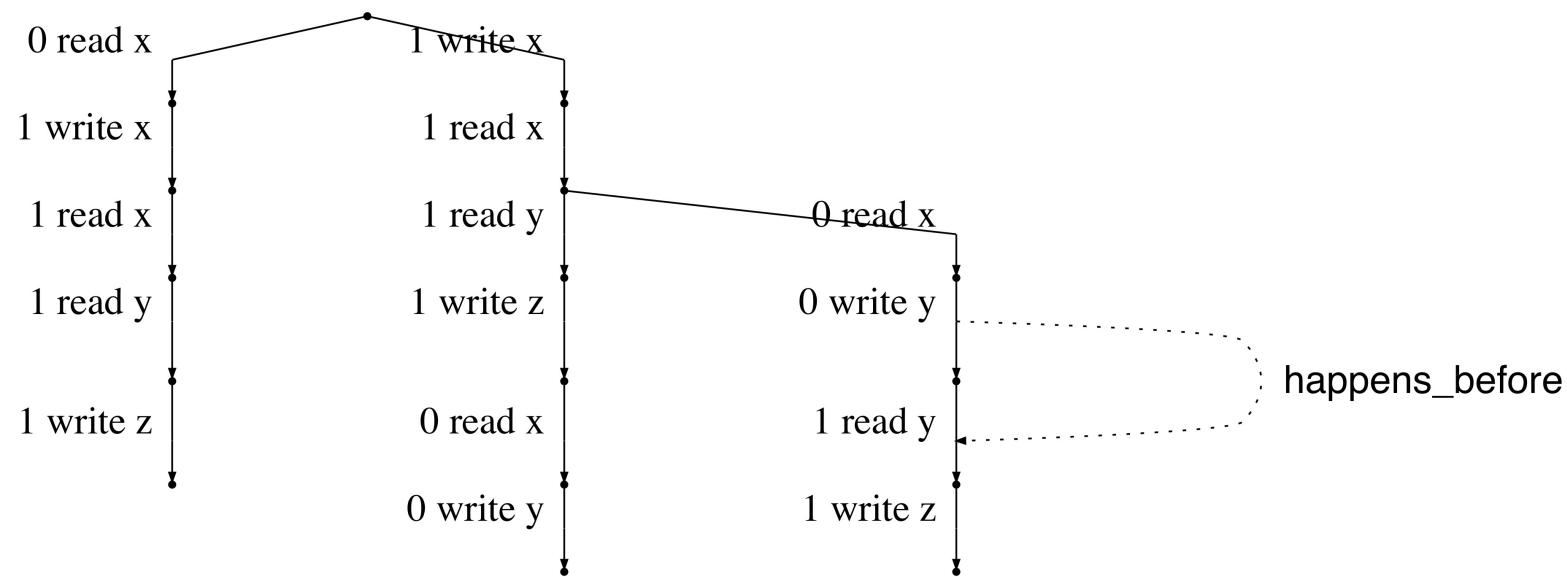
# Partial Order Reduction



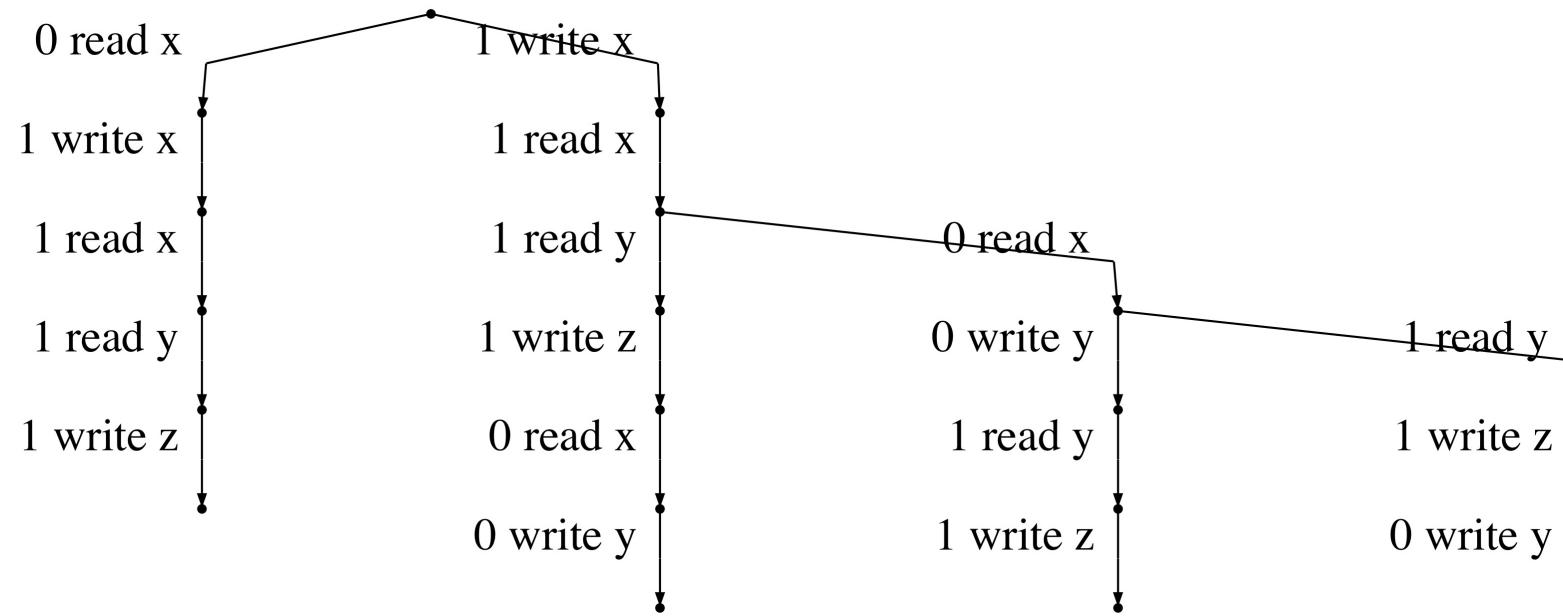
# Partial Order Reduction



# Partial Order Reduction



# Partial Order Reduction



# Partial Order Reduction

## Advantages:

- Complete coverage

## Disadvantages:

- Still infeasible for large state-spaces with many dependencies
- Not incremental

**Thank you!**

# Links

## **Helgrind**

<http://valgrind.org/docs/manual/hg-manual.html>

## **ThreadSanitizer:**

<https://github.com/google/sanitizers/wiki>

## **My Project:**

<https://github.com/s-vde/record-replay>

<https://github.com/s-vde/state-space-explorer>